


# Chapter 3

Arithmetic for Computers



# Arithmetic for Computers

- Operations on integers
  - Addition and subtraction
  - Multiplication and division
  - Dealing with overflow
- Floating-point real numbers
  - Representation and operations

§3.1 Introduction

Chapter 3 — Arithmetic for Computers — RST(c)NTHU

2

§3.2 Addition and Subtraction

## Integer Addition

- Example: 7 + 6

	(0)	(0)	(1)	(1)	(0)	(Carries)
...	0	0	0	1	1	1
...	0	0	0	1	1	0
...	(0)	0	(0)	1	(1)	0
	0	0	1	1	1	0

- Overflow if result out of range
  - Adding +ve and -ve operands, no overflow
  - Adding two +ve operands
    - Overflow if result sign is 1
  - Adding two -ve operands
    - Overflow if result sign is 0

Chapter 3 — Arithmetic for Computers — RST(c)NTHU3

## Integer Subtraction

- Add negation of second operand
- Example: 7 - 6 = 7 + (-6)

+7:	0000	0000	...	0000	0111
-6:	1111	1111	...	1111	1010
+1:	0000	0000	...	0000	0001

- Overflow if result out of range
  - Subtracting two +ve or two -ve operands, no overflow
  - Subtracting +ve from -ve operand
    - Overflow if result sign is 0
  - Subtracting -ve from +ve operand
    - Overflow if result sign is 1

Chapter 3 — Arithmetic for Computers — RST(c)NTHU4

## Dealing with Overflow

- Some languages (e.g., C) ignore overflow
  - Use MIPS `addu`, `addui`, `subu` instructions
- Other languages (e.g., Ada, Fortran) require raising an exception
  - Use MIPS `add`, `addi`, `sub` instructions
  - On overflow, invoke **exception handler**
    - Save PC in exception program counter (EPC) register
    - Jump to predefined handler address
    - `mfc0` (move from coprocessor reg) instruction can retrieve EPC value, to return after corrective action

Chapter 3 — Arithmetic for  
Computers — RST(c)NTHU

5

## Arithmetic for Multimedia

- Graphics and media processing operates on vectors of 8-bit and 16-bit data
  - Use 64-bit adder, with partitioned carry chain
    - Operate on 8×8-bit, 4×16-bit, or 2×32-bit vectors
  - SIMD (single-instruction, multiple-data)
- Saturating operations
  - On overflow, result is largest representable value
    - c.f. 2's-complement modulo arithmetic
  - E.g., clipping in audio, saturation in video

Chapter 3 — Arithmetic for  
Computers — RST(c)NTHU

6

§3.3 Multiplication

# Multiplication

- Start with long-multiplication approach

multiplicand

multiplier

product

1000

× 1001

-----

1000

0000

0000

1000

-----

1001000

Length of product is the sum of operand lengths

Multiplicand

Shift left

64 bits

64-bit ALU

Product

Write

64 bits

Control test

Multiplier

Shift right

32 bits

Chapter 3 — Arithmetic for Computers — RST(c)NTHU

7

# Multiplication Hardware

Start

1. Test Multiplier0

Multiplier0 = 1

Multiplier0 = 0

1a. Add multiplicand to product and place the result in Product register

2. Shift the Multiplicand register left 1 bit

3. Shift the Multiplier register right 1 bit

32nd repetition?

No: < 32 repetitions

Yes: 32 repetitions

Done

Multiplicand

Shift left

64 bits

64-bit ALU

Product

Write

64 bits

Control test

Multiplier

Shift right

32 bits

Initially 0

Chapter 3 — Arithmetic for Computers — RST(c)NTHU

8

## Optimized Multiplier

- Perform steps in parallel: add/shift

- One cycle per partial-product addition
  - That's ok, if frequency of multiplications is low

Chapter 3 — Arithmetic for Computers — RST(c)NTHU

9

## Faster Multiplier

- Uses multiple adders
  - Cost/performance tradeoff

- Can be pipelined
  - Several multiplication performed in parallel

Chapter 3 — Arithmetic for Computers — RST(c)NTHU

10

# MIPS Multiplication

- Two 32-bit registers for product
  - HI: most-significant 32 bits
  - LO: least-significant 32-bits
- Instructions
  - `mult rs, rt` / `multu rs, rt`
    - 64-bit product in HI/LO
  - `mfhi rd` / `mflo rd`
    - Move from HI/LO to rd
    - Can test HI value to see if product overflows 32 bits
  - `mul rd, rs, rt`
    - Least-significant 32 bits of product → rd

Chapter 3 — Arithmetic for Computers — RST(c)NTHU

11

# Division

quotient

dividend

divisor

remainder

1001

1000 ) 1001010

-1000

10

101

1010

-1000

10

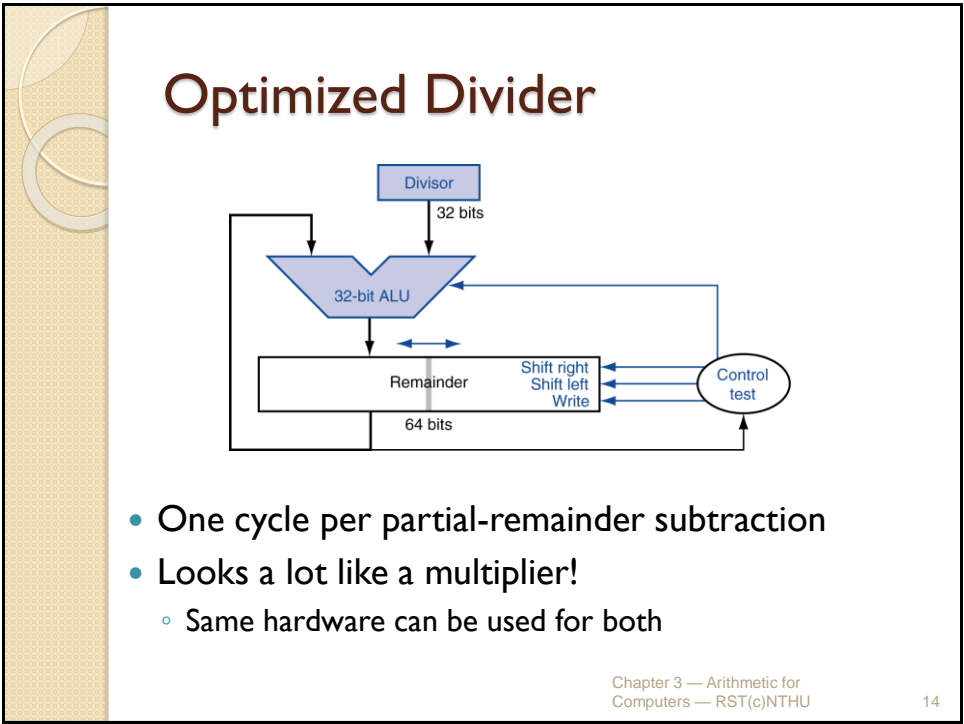
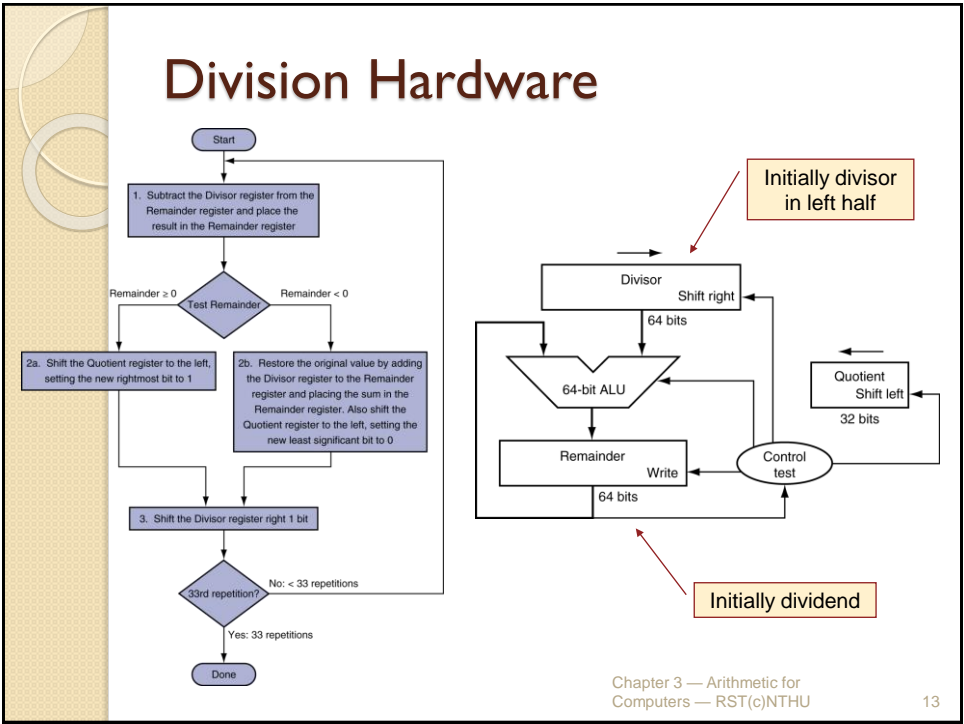
$n$ -bit operands yield  $n$ -bit quotient and remainder

- Check for 0 divisor
- Long division approach
  - If divisor  $\leq$  dividend bits
    - 1 bit in quotient, subtract
  - Otherwise
    - 0 bit in quotient, bring down next dividend bit
- Restoring division
  - Do the subtract, and if remainder goes  $< 0$ , add divisor back
- Signed division
  - Divide using absolute values
  - Adjust sign of quotient and remainder as required

§3.4 Division

Chapter 3 — Arithmetic for Computers — RST(c)NTHU

12



## Faster Division

- Can't use parallel hardware as in multiplier
  - Subtraction is conditional on sign of remainder
- Faster dividers (e.g. SRT division) generate multiple quotient bits per step
  - Still require multiple steps

Chapter 3 — Arithmetic for  
Computers — RST(c)NTHU

15

## MIPS Division

- Use HI/LO registers for result
  - HI: 32-bit remainder
  - LO: 32-bit quotient
- Instructions
  - `div rs, rt` / `divu rs, rt`
  - No overflow or divide-by-0 checking
    - Software must perform checks if required
  - Use `mfhi`, `mflo` to access result

Chapter 3 — Arithmetic for  
Computers — RST(c)NTHU

16



§3.5 Floating Point

## Floating Point

- Representation for non-integral numbers
  - Including very small and very large numbers
- Like scientific notation
  - $-2.34 \times 10^{56}$  ← normalized
  - $+0.002 \times 10^{-4}$  ← not normalized
  - $+987.02 \times 10^9$  ← not normalized
- In binary
  - $\pm 1.xxxxxxx_2 \times 2^{yyyy}$
- Types float and double in C

Chapter 3 — Arithmetic for Computers — RST(c)NTHU17

## Floating Point Standard

- Defined by IEEE Std 754-1985
- Developed in response to divergence of representations
  - Portability issues for scientific code
- Now almost universally adopted
- Two representations
  - Single precision (32-bit) (float)
  - Double precision (64-bit) (double)

Chapter 3 — Arithmetic for Computers — RST(c)NTHU18

## IEEE Floating-Point Format

single: 8 bits  
double: 11 bits

single: 23 bits  
double: 52 bits

S	Exponent	Fraction
---	----------	----------

$$x = (-1)^S \times (1 + \text{Fraction}) \times 2^{(\text{Exponent} - \text{Bias})}$$

- S: sign bit (0  $\Rightarrow$  non-negative, 1  $\Rightarrow$  negative)
- Normalize **significand**:  $1.0 \leq |\text{significand}| < 2.0$ 
  - Always has a leading pre-binary-point 1 bit, so no need to represent it explicitly (hidden bit)
  - Significand is Fraction with the “1.” restored
- Exponent: excess representation: actual exponent + Bias
  - Ensures exponent is unsigned
  - Single: Bias = 127(01111111); Double: Bias = 1203(0111111111)

Chapter 3 — Arithmetic for  
Computers — RST(c)NTHU


19

## Single-Precision Range

- Exponents 00000000 and 11111111 reserved
- Smallest value
  - Exponent: 00000001  
 $\Rightarrow$  actual exponent =  $1 - 127 = -126$
  - Fraction: 000...00  $\Rightarrow$  significand = 1.0
  - $\pm 1.0 \times 2^{-126} \approx \pm 1.2 \times 10^{-38} \rightarrow 0$
- Largest value
  - exponent: 11111110  
 $\Rightarrow$  actual exponent =  $254 - 127 = +127$
  - Fraction: 111...11  $\Rightarrow$  significand  $\approx 2.0$
  - $\pm 2.0 \times 2^{+127} \approx \pm 3.4 \times 10^{+38} \rightarrow \infty$

Chapter 3 — Arithmetic for  
Computers — RST(c)NTHU

20




## Double-Precision Range

- Exponents 0000...00 and 1111...11 reserved
- Smallest value
  - Exponent: 0000000001  
⇒ actual exponent = 1 − 1023 = −1022
  - Fraction: 000...00 ⇒ significand = 1.0
  - $\pm 1.0 \times 2^{-1022} \approx \pm 2.2 \times 10^{-308} \rightarrow 0$
- Largest value
  - Exponent: 1111111110  
⇒ actual exponent = 2046 − 1023 = +1023
  - Fraction: 111...11 ⇒ significand  $\approx 2.0$
  - $\pm 2.0 \times 2^{+1023} \approx \pm 1.8 \times 10^{+308} \rightarrow \infty$

Chapter 3 — Arithmetic for  
Computers — RST(c)NTHU

21




## Floating-Point Precision

- Relative precision
  - all fraction bits are significant
  - Single: approx  $2^{-23}$ 
    - Equivalent to  $23 \times \log_{10} 2 \approx 23 \times 0.3 \approx 6$  decimal digits of precision
  - Double: approx  $2^{-52}$ 
    - Equivalent to  $52 \times \log_{10} 2 \approx 52 \times 0.3 \approx 16$  decimal digits of precision

Chapter 3 — Arithmetic for  
Computers — RST(c)NTHU

22




## Floating-Point Example

- Represent  $-0.75$ 
  - $-0.75 = (-1)^1 \times 1.1_2 \times 2^{-1}$
  - $S = 1$
  - Fraction =  $1000\dots00_2$
  - Exponent =  $-1 + \text{Bias}$ 
    - Single:  $-1 + 127 = 126 = 0111110_2$
    - Double:  $-1 + 1023 = 1022 = 011111110_2$
- Single: 101111101000...00
- Double: 10111111101000...00

Chapter 3 — Arithmetic for  
Computers — RST(c)NTHU

23



## Floating-Point Example

- What number is represented by the single-precision float  
11000000101000...00
  - $S = 1$
  - Fraction =  $01000\dots00_2$
  - Exponent = 1000000 $_2 = 129$
- $x = (-1)^1 \times (1 + .01_2) \times 2^{(129 - 127)}$ 
$$= (-1) \times 1.25 \times 2^2$$
$$= -5.0$$

Chapter 3 — Arithmetic for  
Computers — RST(c)NTHU

24

## Denormalized Numbers

- Exponent = 000...0  $\Rightarrow$  hidden bit is 0
$$x = (-1)^S \times (0 + .\text{Fraction}) \times 2^{-\text{Bias}}$$
- Smaller than normal numbers
  - allow for gradual underflow, with diminishing precision
- Denormal with fraction = 000...0
$$x = (-1)^S \times (0 + .0) \times 2^{-\text{Bias}} = \pm 0.0$$

Two representations of 0.0!

Chapter 3 — Arithmetic for Computers — RST(c)NTHU

25

## Infinities and NaNs

- Exponent = 111...1, Fraction = 000...0
  - $\pm$ Infinity
  - Can be used in subsequent calculations, avoiding need for overflow check
- Exponent = 111...1, Fraction  $\neq$  000...0
  - Not-a-Number (NaN)
  - Indicates illegal or undefined result
    - e.g., 0.0 / 0.0
  - Can be used in subsequent calculations

Chapter 3 — Arithmetic for Computers — RST(c)NTHU

26

## Floating-Point Addition

Consider a 4-digit decimal example

$$9.999 \times 10^1 + 1.610 \times 10^{-1}$$

1. Align decimal points
  - Shift number with **smaller** exponent
  - $9.999 \times 10^1 + 0.016 \times 10^1$
2. Add significands
  - $9.999 \times 10^1 + 0.016 \times 10^1 = 10.015 \times 10^1$
3. Normalize result & check for over/underflow
  - $1.0015 \times 10^2$
4. Round and renormalize if necessary
  - $1.002 \times 10^2$

Chapter 3 — Arithmetic for  
Computers — RST(c)NTHU

27

## Floating-Point Addition

Now consider a 4-digit binary example

$$1.000_2 \times 2^{-1} + -1.110_2 \times 2^{-2} \quad (0.5 + -0.4375)$$

1. Align binary points
  - Shift number with smaller exponent
  - $1.000_2 \times 2^{-1} + -0.111_2 \times 2^{-1}$
2. Add significands
  - $1.000_2 \times 2^{-1} + -0.111_2 \times 2^{-1} = 0.001_2 \times 2^{-1}$
3. Normalize result & check for over/underflow
  - $1.000_2 \times 2^{-4}$ , with no over/underflow
4. Round and renormalize if necessary
  - $1.000_2 \times 2^{-4}$  (no change) = 0.0625

Chapter 3 — Arithmetic for  
Computers — RST(c)NTHU

28

## FP Adder Hardware

- Much more complex than integer adder
- Doing it in one clock cycle would take too long
  - Much longer than integer operations
  - Slower clock would penalize all instructions
- FP adder usually takes several cycles
  - Can be pipelined

Chapter 3 — Arithmetic for Computers — RST(c)NTHU29

## FP Adder Hardware

The diagram illustrates the hardware components and data flow for floating-point addition. It starts with two input registers, each containing Sign, Exponent, and Fraction. The exponents are compared in a 'Small ALU' to determine the 'Exponent difference'. This difference is used by a 'Control' unit to manage shifts. The smaller number is shifted right ('Shift right') until the exponents are equal. The fractions are then added in a 'Big ALU'. The result is then normalized ('Shift left or right') and rounded ('Rounding hardware') to produce the final Sign, Exponent, and Fraction.

Annotations on the right side of the diagram:

- Step 1:** Compare exponents
- Step 2:** Add (includes Shift smaller number right)
- Step 3:** Normalize
- Step 4:** Round

Chapter 3 — Arithmetic for Computers — RST(c)NTHU30

## FP Arithmetic Hardware

- FP multiplier is of similar complexity to FP adder
  - But uses a multiplier for significands instead of an adder
- FP arithmetic hardware usually does
  - Addition, subtraction, multiplication, division, reciprocal, square-root
  - $FP \leftrightarrow$  integer conversion
- Operations usually takes several cycles
  - Can be pipelined

Chapter 3 — Arithmetic for  
Computers — RST(c)NTHU

33

## FP Instructions in MIPS

- FP hardware is coprocessor I
  - Adjunct processor that extends the ISA
- Separate FP registers
  - 32 single-precision:  $\$f0, \$f1, \dots, \$f31$
  - Paired for double-precision:  $\$f0/\$f1, \$f2/\$f3, \dots$ 
    - Release 2 of MIPS ISA supports  $32 \times 64$ -bit FP reg's
- FP instructions operate only on FP registers
  - Programs generally don't do integer ops on FP data, or vice versa
  - More registers with minimal code-size impact
- FP load and store instructions
  - $lwc1, ldc1, swc1, sdc1$ 
    - e.g.,  $ldc1 \$f8, 32(\$sp)$

Chapter 3 — Arithmetic for  
Computers — RST(c)NTHU

34



## FP Instructions in MIPS

- Single-precision arithmetic
  - `add.s, sub.s, mul.s, div.s`
    - e.g., `add.s $f0, $f1, $f6`
- Double-precision arithmetic
  - `add.d, sub.d, mul.d, div.d`
    - e.g., `mul.d $f4, $f4, $f6`
- Single- and double-precision comparison
  - `c.xx.s, c.xx.d` (`xx` is `eq, lt, le, ...`)
  - Sets or clears FP condition-code bit
    - e.g., `c.lt.s $f3, $f4`
- Branch on FP condition code true or false
  - `bc1t, bc1f`
    - e.g., `bc1t TargetLabel`

Chapter 3 — Arithmetic for  
Computers — RST(c)NTHU

35

## FP Example: °F to °C

- C code:
 

```
float f2c (float fahr) {
    return ((5.0/9.0)*(fahr - 32.0));
}
```

  - `fahr` in `$f12`, result in `$f0`, literals in global memory space
- Compiled MIPS code:
 

```
f2c: lwc1    $f16, const5($gp)
      lwc2    $f18, const9($gp)
      div.s   $f16, $f16, $f18
      lwc1    $f18, const32($gp)
      sub.s   $f18, $f12, $f18
      mul.s   $f0, $f16, $f18
      jr      $ra
```

Chapter 3 — Arithmetic for  
Computers — RST(c)NTHU

36

## FP Example:Array Multiplication

- $X = X + Y \times Z$ 
  - All  $32 \times 32$  matrices, 64-bit double-precision elements
- C code:

```
void mm (double x[][],
         double y[][], double z[][]) {
    int i, j, k;
    for (i = 0; i != 32; i = i + 1)
        for (j = 0; j != 32; j = j + 1)
            for (k = 0; k != 32; k = k + 1)
                x[i][j] = x[i][j]
                    + y[i][k] * z[k][j];
}
```

  - Addresses of x, y, z in \$a0, \$a1, \$a2, and i, j, k in \$s0, \$s1, \$s2

Chapter 3 — Arithmetic for Computers — RST(c)NTHU

37

## FP Example:Array Multiplication

- MIPS code:

li	\$t1, 32	# \$t1 = 32 (row size/loop end)
li	\$s0, 0	# i = 0; initialize 1st for loop
L1:	li	\$s1, 0 # j = 0; restart 2nd for loop
L2:	li	\$s2, 0 # k = 0; restart 3rd for loop
sll	\$t2, \$s0, 5	# \$t2 = i * 32 (size of row of x)
addu	\$t2, \$t2, \$s1	# \$t2 = i * size(row) + j
sll	\$t2, \$t2, 3	# \$t2 = byte offset of [i][j]
addu	\$t2, \$a0, \$t2	# \$t2 = byte address of x[i][j]
l.d	\$f4, 0(\$t2)	# \$f4 = 8 bytes of x[i][j]
L3:	sll	\$t0, \$s2, 5 # \$t0 = k * 32 (size of row of z)
addu	\$t0, \$t0, \$s1	# \$t0 = k * size(row) + j
sll	\$t0, \$t0, 3	# \$t0 = byte offset of [k][j]
addu	\$t0, \$a2, \$t0	# \$t0 = byte address of z[k][j]
l.d	\$f16, 0(\$t0)	# \$f16 = 8 bytes of z[k][j]
...		

Chapter 3 — Arithmetic for Computers — RST(c)NTHU

38

## FP Example: Array Multiplication

...

sll	\$t0, \$s0, 5	# \$t0 = i*32 (size of row of y)
addu	\$t0, \$t0, \$s2	# \$t0 = i*size(row) + k
sll	\$t0, \$t0, 3	# \$t0 = byte offset of [i][k]
addu	\$t0, \$a1, \$t0	# \$t0 = byte address of y[i][k]
l.d	\$f18, 0(\$t0)	# \$f18 = 8 bytes of y[i][k]
mul.d	\$f16, \$f18, \$f16	# \$f16 = y[i][k] * z[k][j]
add.d	\$f4, \$f4, \$f16	# f4=x[i][j] + y[i][k]*z[k][j]
addiu	\$s2, \$s2, 1	# \$k k + 1
bne	\$s2, \$t1, L3	# if (k != 32) go to L3
s.d	\$f4, 0(\$t2)	# x[i][j] = \$f4
addiu	\$s1, \$s1, 1	# \$j = j + 1
bne	\$s1, \$t1, L2	# if (j != 32) go to L2
addiu	\$s0, \$s0, 1	# \$i = i + 1
bne	\$s0, \$t1, L1	# if (i != 32) go to L1

Chapter 3 — Arithmetic for Computers — RST(c)NTHU

39

## Subword Parallellism

- Graphics and audio applications can take advantage of performing simultaneous operations on short vectors
  - Example: 128-bit adder:
    - Sixteen 8-bit adds
    - Eight 16-bit adds
    - Four 32-bit adds
- Also called data-level parallelism, vector parallelism, or Single Instruction, Multiple Data (SIMD)

Chapter 3 — Arithmetic for Computers — 41

§3.6 Parallelism and Computer Arithmetic: Subword Parallellism

§3.7 Real Stuff: Floating Point in the x86

## x86 FP Architecture

- Originally based on 8087 FP coprocessor
  - 8 × 80-bit extended-precision registers
  - Used as a push-down stack
  - Registers indexed from TOS: ST(0), ST(1), ...
- FP values are 32-bit or 64 in memory
  - Converted on load/store of memory operand
  - Integer operands can also be converted on load/store
- Very difficult to generate and optimize code
  - Result: poor FP performance

Chapter 3 — Arithmetic for Computers — RST(c)NTHU

43


## x86 FP Instructions

Data transfer	Arithmetic	Compare	Transcendental
<b>F</b> ILD mem/ST(i)	<b>F</b> I <b>A</b> DD <b>P</b> mem/ST(i)	<b>F</b> I <b>C</b> OMP	FPATAN
<b>F</b> I <b>S</b> TP mem/ST(i)	<b>F</b> I <b>S</b> UB <b>R</b> P mem/ST(i)	<b>F</b> I <b>U</b> COMP	F2XMI
FLDPI	<b>F</b> I <b>M</b> UL <b>P</b> mem/ST(i)	FSTSW AX/mem	FCOS
FLD1	<b>F</b> I <b>D</b> IV <b>R</b> P mem/ST(i)		FPTAN
FLDZ	FSQRT		FPREM
	FABS		FPSIN
	FRNDINT		FYL2X

- Optional variations
  - I**: integer operand
  - P**: pop operand from stack
  - R**: reverse operand order
  - But not all combinations allowed

Chapter 3 — Arithmetic for Computers — RST(c)NTHU

44




## Streaming SIMD Extension 2 (SSE2)

- Adds 4 × 128-bit registers
  - Extended to 8 registers in AMD64/EM64T
- Can be used for multiple FP operands
  - 2 × 64-bit double precision
  - 4 × 32-bit double precision
  - Instructions operate on them simultaneously
    - Single-Instruction Multiple-Data

Chapter 3 — Arithmetic for Computers — RST(c)NTHU

45



## Matrix Multiply

- Unoptimized code:

```
1. void dgemm (int n, double* A, double* B, double* C)
2. {
3.     for (int i = 0; i < n; ++i)
4.         for (int j = 0; j < n; ++j)
5.             {
6.                 double cij = C[i+j*n]; /* cij = C[i][j] */
7.                 for(int k = 0; k < n; k++ )
8.                     cij += A[i+k*n] * B[k+j*n]; /* cij += A[i][k]*B[k][j] */
9.                 C[i+j*n] = cij; /* C[i][j] = cij */
10.            }
11. }
```

§3.8 Going Faster: Subword Parallelism and Matrix Multiply

# Matrix Multiply


- x86 assembly code:

```
1. vmovsd (%r10),%xmm0 # Load 1 element of C into %xmm0
2. mov %rsi,%rcx        # register %rcx = %rsi
3. xor %eax,%eax        # register %eax = 0
4. vmovsd (%rcx),%xmm1  # Load 1 element of B into %xmm1
5. add %r9,%rcx         # register %rcx = %rcx + %r9
6. vmulsd (%r8,%rax,8),%xmm1,%xmm1 # Multiply %xmm1,
element of A
7. add $0x1,%rax        # register %rax = %rax + 1
8. cmp %eax,%edi        # compare %eax to %edi
9. vaddsd %xmm1,%xmm0,%xmm0 # Add %xmm1, %xmm0
10. jg 30 <dgemm+0x30>  # jump if %eax > %edi
11. add $0x1,%r11d      # register %r11 = %r11 + 1
12. vmovsd %xmm0, (%r10) # Store %xmm0 into C element
```

# Matrix Multiply


- Optimized C code:

```
1. #include <x86intrin.h>
2. void dgemm (int n, double* A, double* B, double* C)
3. {
4.     for ( int i = 0; i < n; i+=4 )
5.         for ( int j = 0; j < n; j++ ) {
6.             __m256d c0 = __mm256_load_pd(C+i+j*n); /* c0 = C[i][j] */
7.             for( int k = 0; k < n; k++ )
8.                 c0 = __mm256_add_pd(c0, /* c0 += A[i][k]*B[k][j] */
9.                                     __mm256_mul_pd(__mm256_load_pd(A+i+k*n),
10.                                                     __mm256_broadcast_sd(B+k+j*n)));
11.             __mm256_store_pd(C+i+j*n, c0); /* C[i][j] = c0 */
12.         }
13. }
```



## Matrix Multiply

- Optimized x86 assembly code:
  1. vmovapd (%r11),%ymm0 # Load 4 elements of C into %ymm0
  2. mov %rbx,%rcx # register %rcx = %rbx
  3. xor %eax,%eax # register %eax = 0
  4. vbroadcastsd (%rax,%r8,1),%ymm1 # Make 4 copies of B element
  5. add \$0x8,%rax # register %rax = %rax + 8
  6. vmulpd (%rcx),%ymm1,%ymm1 # Parallel mul %ymm1, 4 A elements
  7. add %r9,%rcx # register %rcx = %rcx + %r9
  8. cmp %r10,%rax # compare %r10 to %rax
  9. vaddpd %ymm1,%ymm0,%ymm0 # Parallel add %ymm1, %ymm0
  10. jne 50 <dgemm+0x50> # jump if not %r10 != %rax
  11. add \$0x1,%esi # register % esi = % esi + 1
  12. vmovapd %ymm0, (%r11) # Store %ymm0 into 4 C elements



## Right Shift and Division

- Left shift by  $i$  places multiplies an integer by  $2^i$
- Right shift divides by  $2^i$ ?
  - Only for unsigned integers
- For signed integers
  - Arithmetic right shift: replicate the sign bit
  - e.g.,  $-5 / 4$ 
    - $|||||011_2 \gg 2 = |||||10_2 = -2$
    - Rounds toward  $-\infty$
  - c.f.  $|||||011_2 \ggg 2 = 00|||||10_2 = +62$

§3.8 Fallacies and Pitfalls

Chapter 3 — Arithmetic for Computers — RST(c)NTHU

50

§3.8 Parallelism and Computer Arithmetic: Associativity

## Associativity

- Parallel programs may interleave operations in unexpected orders
  - Assumptions of associativity may fail

		$(x+y)+z$	$x+(y+z)$
x	-1.50E+38	0.00E+00	-1.50E+38
y	1.50E+38		1.50E+38
z	1.0	1.0	
		1.00E+00	0.00E+00

- Need to validate parallel programs under varying degrees of parallelism

Chapter 3 — Arithmetic for Computers — RST(c)NTHU

51

## Who Cares About FP Accuracy?

- Important for scientific code
  - But for everyday consumer use?
    - “My bank balance is out by 0.0002¢!” ☹
- The Intel Pentium FDIV bug
  - The market expects accuracy
  - See Colwell, *The Pentium Chronicles*

Chapter 3 — Arithmetic for Computers — RST(c)NTHU

52



## Concluding Remarks 1/2

- Bits have no inherent meaning
  - Interpretation depends on the instructions applied
- Computer representations of numbers
  - Finite range and precision
  - Need to account for this in programs

## Concluding Remarks 2/2

- ISAs support arithmetic
  - Signed and unsigned integers
  - Floating-point approximation to reals
- Bounded range and precision
  - Operations can overflow and underflow
- MIPS ISA
  - Core instructions: 54 most frequently used
    - 100% of SPECINT, 97% of SPECFP
  - Other instructions: less frequent