



Chapter 4 The Processor



Introduction

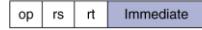
§4.1 Introduction

- CPU performance factors
 - Instruction count
 - Determined by ISA and compiler
 - CPI and Cycle time
 - Determined by CPU hardware
- We will examine two MIPS implementations
 - A simplified version
 - A more realistic pipelined version
- Simple subset, shows most aspects
 - Memory reference: lw, sw
 - Arithmetic/logical: add, sub, and, or, slt
 - Control transfer: beq, j



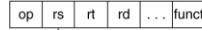
What Hardware to Implement?

1. Immediate addressing



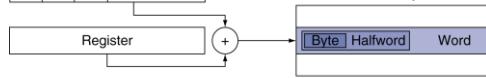
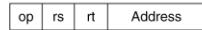
`addi $s0,$s1,100`

2. Register addressing



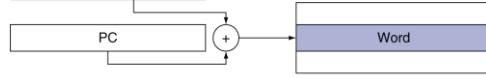
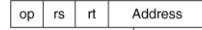
`add $s0,$s1,$s3`

3. Base addressing



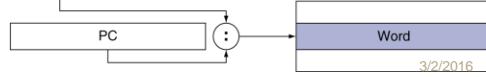
`lw $s0,4($s1)`

4. PC-relative addressing



`beq $s0,$s1,L1`

5. Pseudodirect addressing



`J L2`

3/2/2016 Chapter 4-The Processor- RST@NTHU

32-bit MIPS Instruction Set

	31	26	21	16	11	6	0				
R-type	op	rs	rt	rd	shamt	funct					
	6 bits	5 bits	5 bits	5 bits	5 bits	6 bits					
I-type	op	rs	rt	immediate							
	6 bits	5 bits	5 bits	16 bits							
J-type	op	target address									
	6 bits	26 bits									

The different fields are:

op	operation of the instruction
rs, rt, rd	source and destination register
shamt	shift amount
funct	selects variant of the “op” field
Immediate	address / constant
target address	target address of jump

A MIPS Subset

- R-Type:

- add rd, rs, rt
- sub rd, rs, rt
- and rd, rs, rt
- or rd, rs, rt
- slt rd, rs, rt

- Load/Store:

- lw rt,rs,imm16
- sw rt,rs,imm16

- Imm operand:

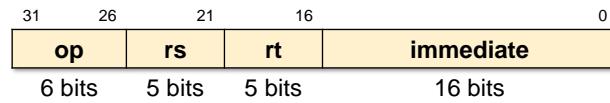
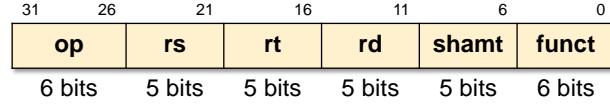
- addi rt,rs,imm16

- Branch:

- beq rs,rt,imm16

- Jump:

- j target



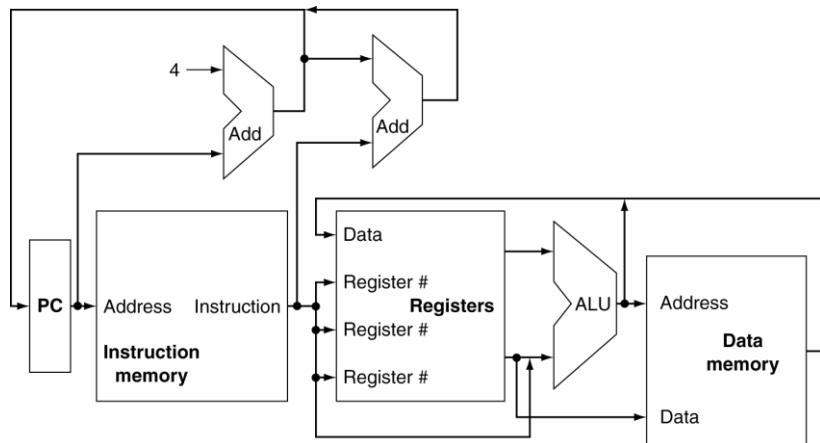
3/2/2016 Chapter 4-The Processor- RST@NTHU

Instruction Execution

- PC → instruction memory, fetch instruction
- Register numbers → register file, read registers
- Depending on instruction class
 - Use ALU to calculate
 - Arithmetic result
 - Memory address for load/store
 - Branch target address
 - Access data memory for load/store
 - PC ← target address or PC + 4

3/2/2016 Chapter 4-The Processor- RST@NTHU

CPU Overview

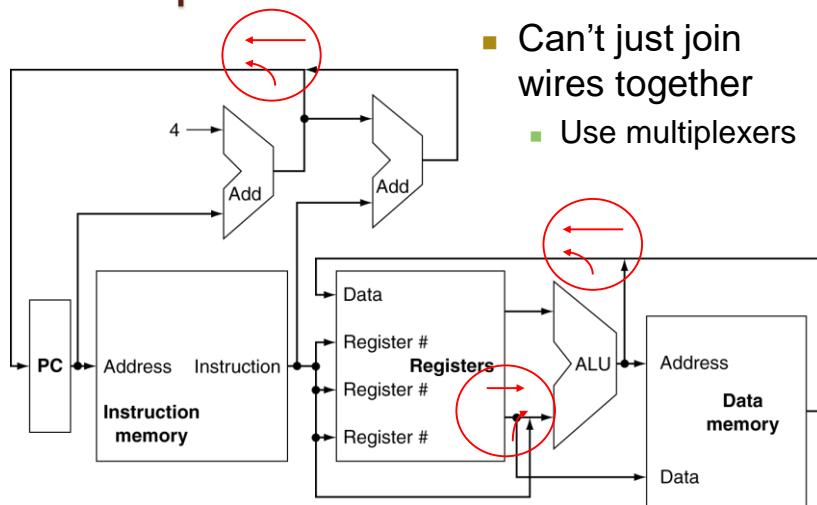


302

3/2/2016 Chapter 4-The Processor- RST@NTHU

Multiplexers

- Can't just join wires together
 - Use multiplexers

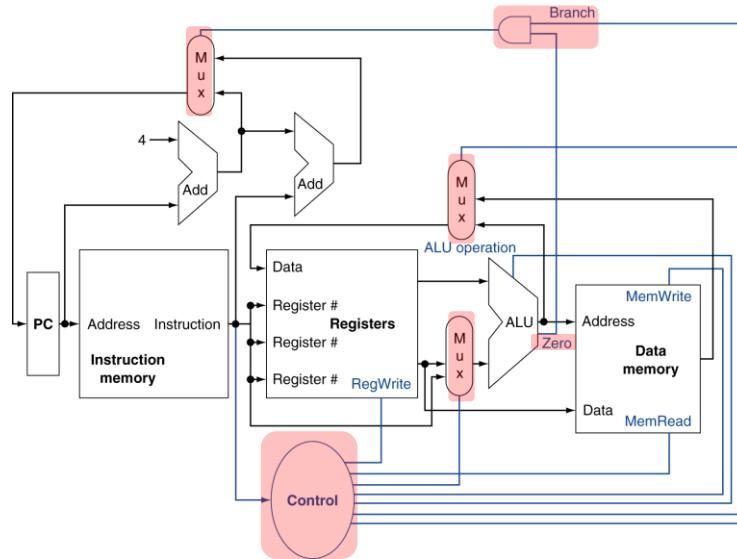


3/2/2016

Chapter 4-The Processor- RST@NTHU



Control



3/2/2016 Chapter 4-The Processor- RST@NTHU



Logic Design Basics

- Information encoded in binary
 - Low voltage = 0, High voltage = 1
 - One wire per bit
 - Multi-bit data encoded on multi-wire buses
- Combinational element
 - Operate on data
 - Output is a function of input
- State (sequential) elements
 - Store information

§4.2 Logic Design Conventions

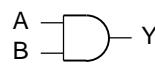
3/2/2016 Chapter 4-The Processor- RST@NTHU



Combinational Elements

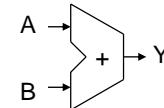
- AND-gate

- $Y = A \& B$



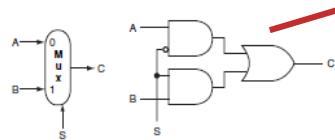
- Adder

- $Y = A + B$



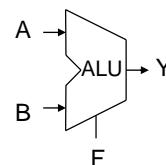
- Multiplexer

- $Y = S ? A : B$



- Arithmetic/Logic Unit

- $Y = F(A, B)$



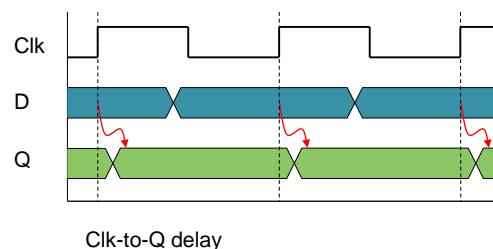
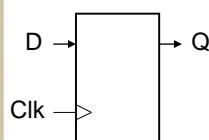
3/2/2016 Chapter 4-The Processor- RST@NTHU



Sequential Elements

- Register: stores data in a circuit

 - Uses a clock signal to determine when to update the stored value
 - Rising edge-triggered: update when Clk changes from 0 to 1

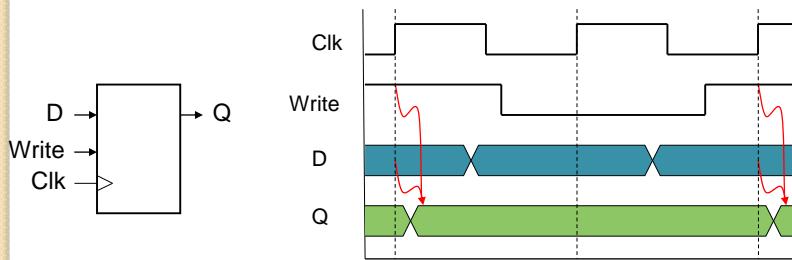


305

3/2/2016 Chapter 4-The Processor- RST@NTHU

Sequential Elements

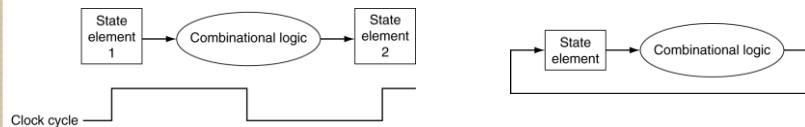
- Register with write control
 - Only updates on clock edge when write control input is 1
 - Used when stored value is required later



3/2/2016 Chapter 4-The Processor- RST@NTHU

Clocking Methodology

- Combinational logic transforms data during clock cycles
 - Between clock edges
 - Input from state elements, output to state element
 - Longest delay determines clock period



306

3/2/2016 Chapter 4-The Processor- RST@NTHU

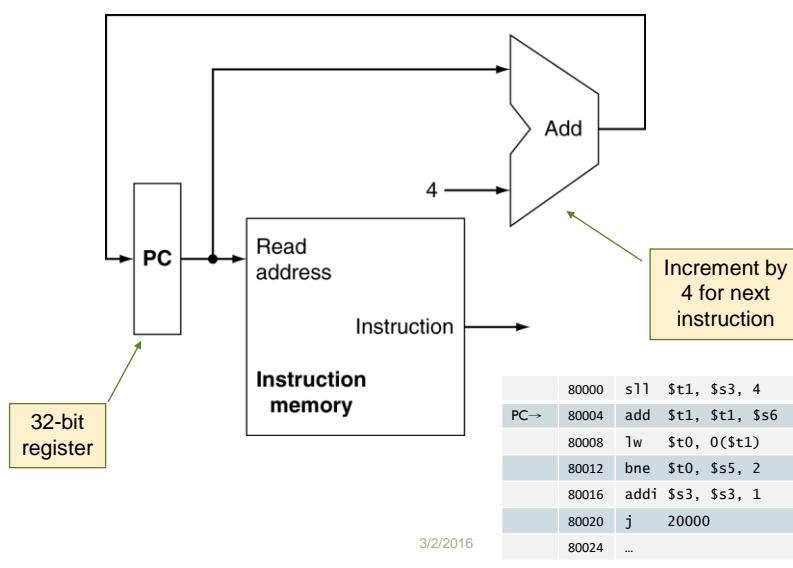
Building a Datapath

- Datapath
 - Elements that process data and addresses in the CPU
 - Registers, ALUs, mux's, memories, ...
- We will build a MIPS datapath incrementally
 - Refining the overview design

239

3/2/2016 Chapter 4-The Processor- RST@NTHU

Instruction Fetch



241

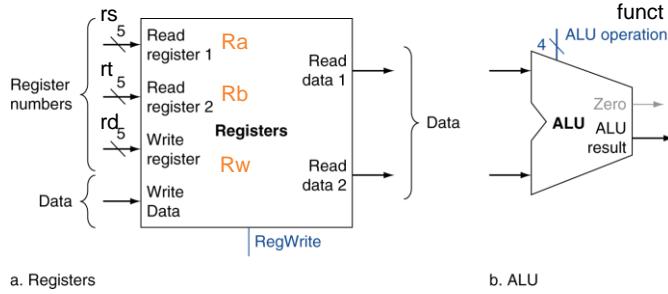
3/2/2016



R-Format Instructions

- Read two register operands
- Perform arithmetic/logical operation
- Write register result
- $R[rd] \leftarrow R[rs] \text{ op } R[rt]$ Ex: add rd, rs, rt
 - Ra, Rb, Rw come from inst.'s rs, rt, and rd fields
 - ALU and RegWrite: control logic after decode

op	rs	rt	rd	shamt	funct
----	----	----	----	-------	-------



241

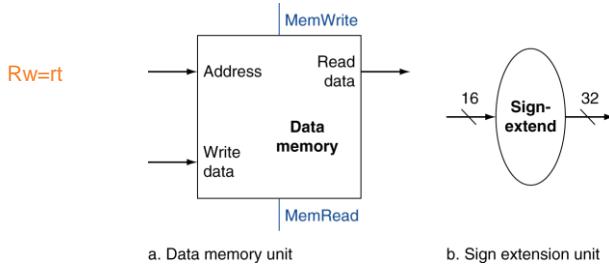
3/2/2016 Chapter 4-The Processor- RST@NTHU



Load/Store Instructions

- Read register operands
- Calculate address using 16-bit offset
 - Use ALU, but sign-extend offset
- Load: Read memory and update register
- Store: Write register value to memory
- $R[rt] \leftarrow \text{Mem}[R[rs] + \text{SignExt}[\text{imm16}]]$ Ex: lw rt,rs,imm16

op	rs	rt	constant or address
----	----	----	---------------------



243

3/2/2016 Chapter 4-The Processor- RST@NTHU



Branch Instructions

op | rs | rt | constant or address

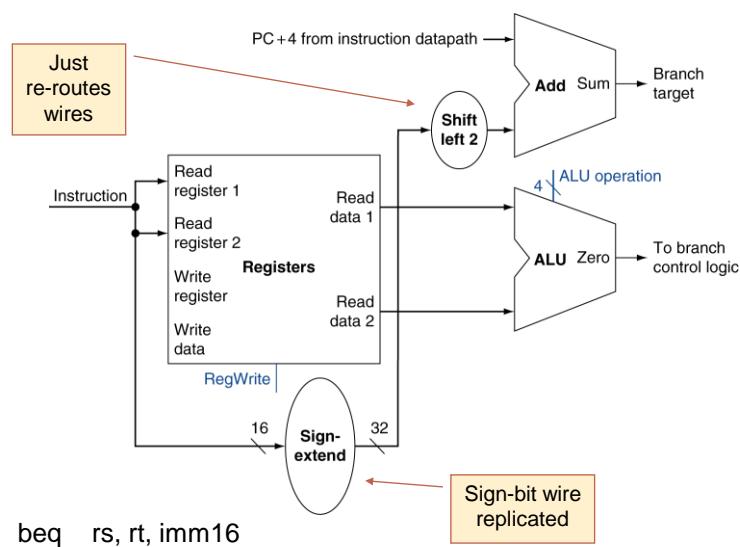
- Read register operands
- Compare operands
 - Use ALU, subtract and check Zero output
- Calculate target address
 - Sign-extend displacement
 - Shift left 2 places (word displacement)
 - Add to PC + 4
 - Already calculated by instruction fetch

242

3/2/2016 Chapter 4-The Processor- RST@NTU



Branch Instructions



beq rs, rt, imm16

op | rs | rt | constant or address

3/2/2016 Chapter 4-The Processor- RST@NTU



Composing the Elements

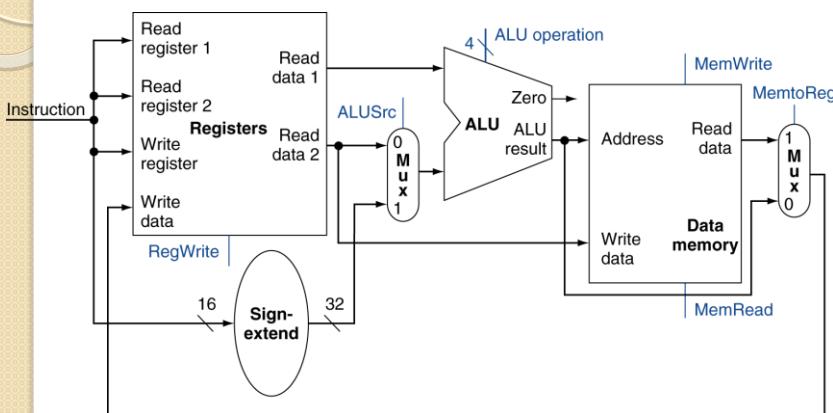
- First-cut data path does one instruction in one clock cycle
 - Each datapath element can only do one function at a time
 - Hence, we need separate instruction and data memories
- Use multiplexers where alternate data sources are used for different instructions

313

3/2/2016 Chapter 4-The Processor- RST@NTHU



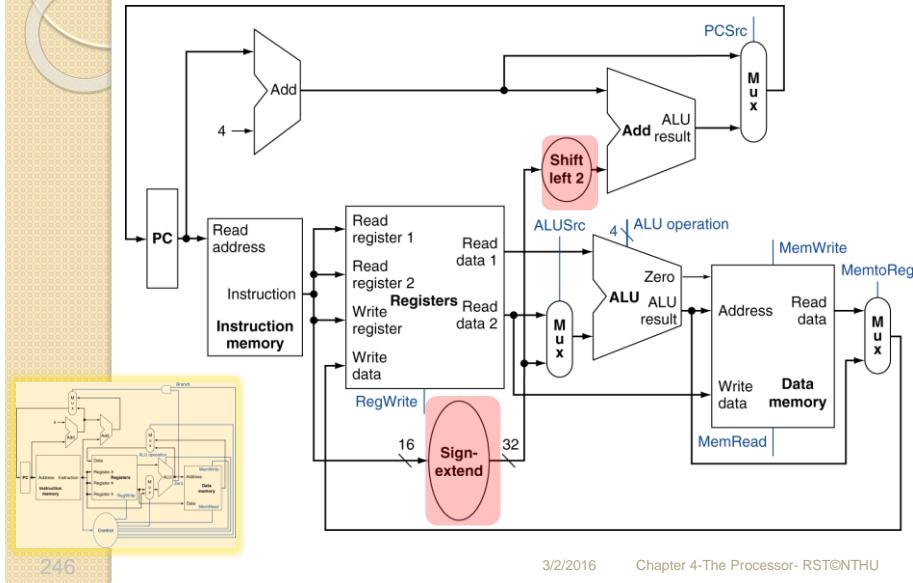
R-Type/Load/Store Datapath



246

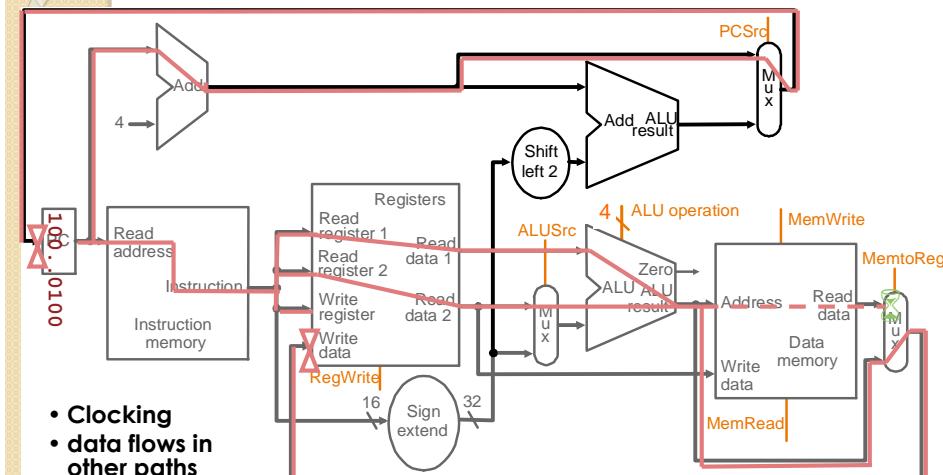
3/2/2016 Chapter 4-The Processor- RST@NTHU

Full Datapath

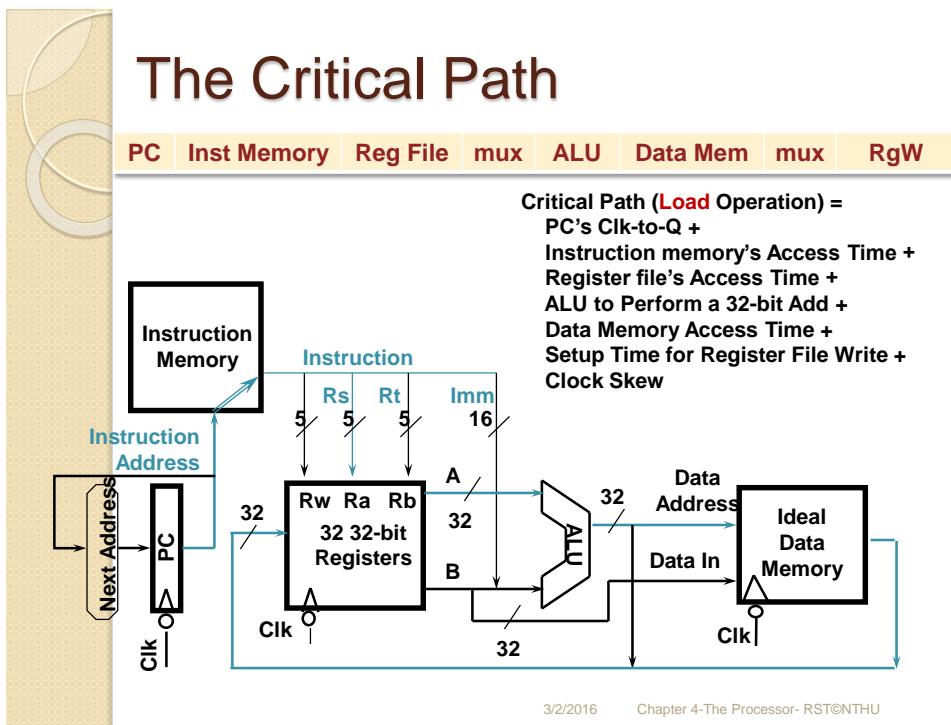
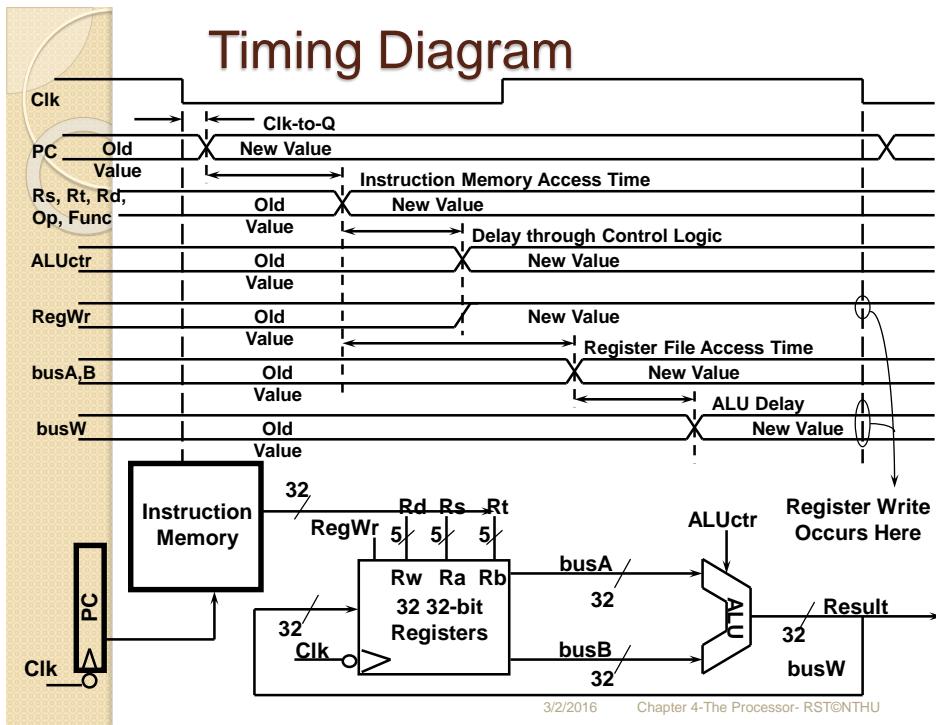


3/2/2016 Chapter 4-The Processor- RST@NTHU

Data Flow for add rd, rs, rt



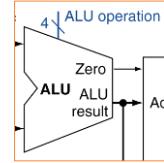
3/2/2016 Chapter 4-The Processor- RST@NTHU



§4.4 A Simple Implementation Scheme

ALU Control

000000 | rs | rt | rd | shamt | funct



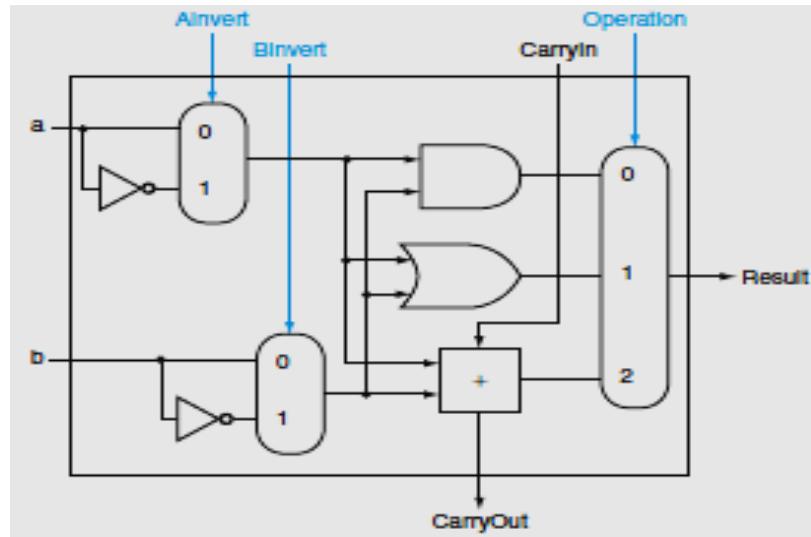
- ALU used for
 - Load/Store: $F = \text{add}$
 - Branch: $F = \text{subtract}$
 - R-type: F depends on funct field

ALU control	Function
0000	AND
0001	OR
0010	add
0110	subtract
0111	set-on-less-than
1100	NOR

247

3/2/2016 Chapter 4-The Processor- RST@NTHU

A 1-bit ALU: and, or, add



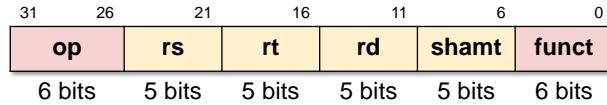
C-32

3/2/2016 Chapter 4-The Processor- RST@NTHU



ALU Control

- Assume 2-bit ALUOp derived from opcode
 - Combinational logic derives ALU control



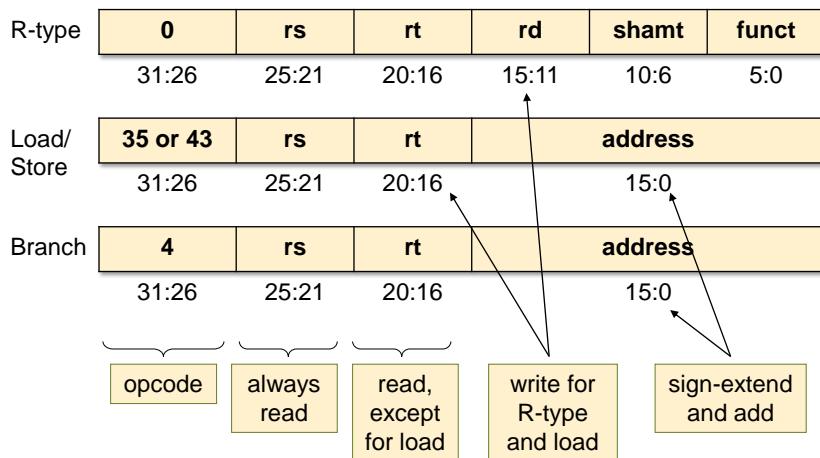
opcode	ALUOp	Operation	funct	ALU function	ALU control
lw	00	load word	xxxxxx	add	0010
sw	00	store word	xxxxxx	add	0010
beq	01	branch equal	xxxxxx	subtract	0110
R-type	10	add	100000	add	0010
		subtract	100010	subtract	0110
		AND	100100	AND	0000
		OR	100101	OR	0001
		set-on-less-than	101010	set-on-less-than	0111

3/2/2016 Chapter 4-The Processor- RST@NTHU



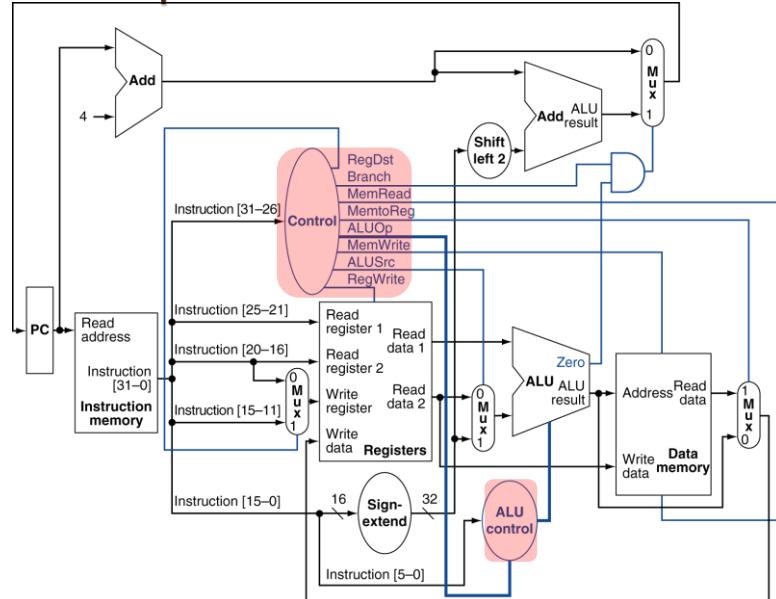
The Main Control Unit

- Control signals derived from instruction



3/2/2016 Chapter 4-The Processor- RST@NTHU

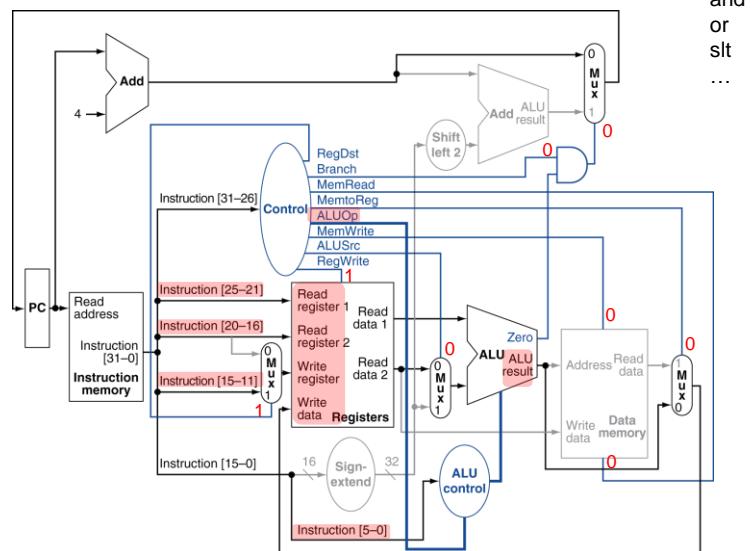
Datapath With Control



253

3/2/2016 Chapter 4-The Processor- RST@NTHU

R-Type Instruction

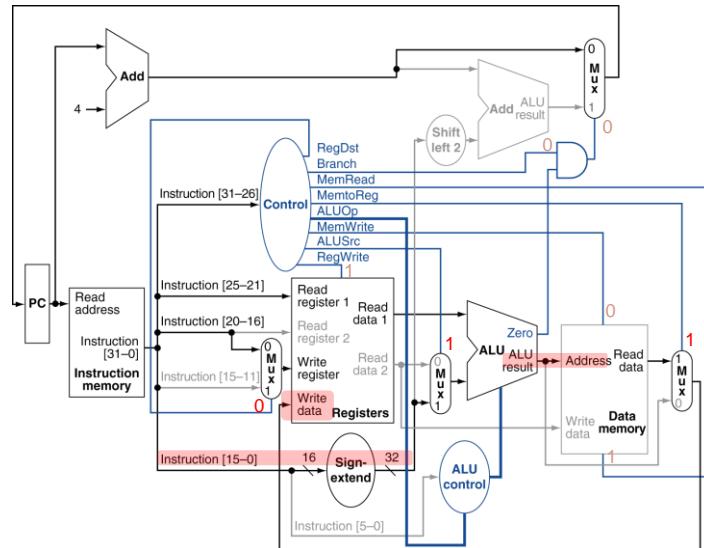


254

add
sub
and
or
slt
...

3/2/2016 Chapter 4-The Processor- RST@NTHU

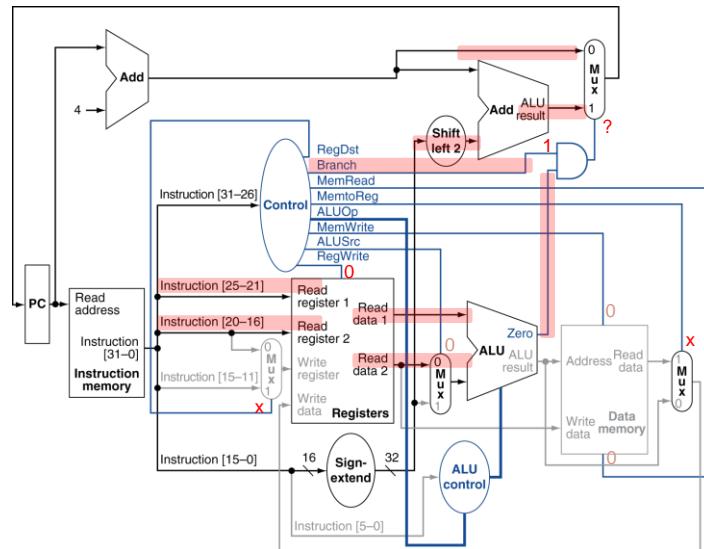
Load Instruction lw \$s1, 4(\$t0)



255

3/2/2016 Chapter 4-The Processor- RST@NTHU

Branch-on-Equal Instruction



256

3/2/2016 Chapter 4-The Processor- RST@NTHU

Implementing Jumps

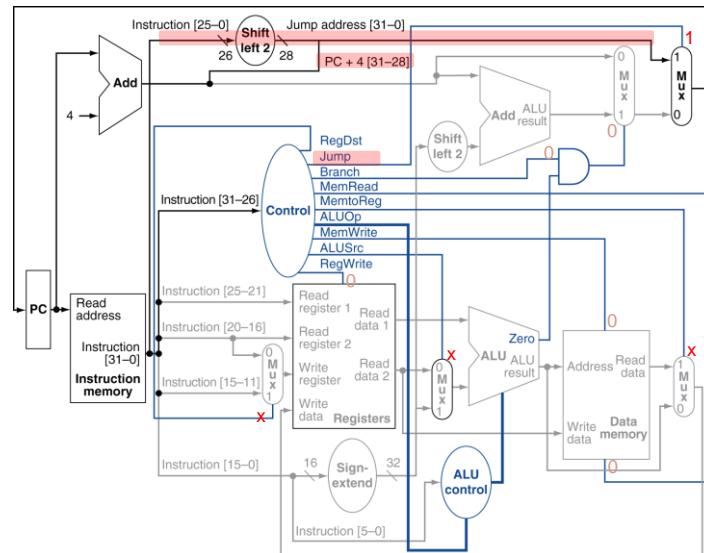


- Jump uses word address
- Update PC with concatenation of
 - Top 4 bits of old PC $\text{PC}_{31\ldots 28}$
 - 26-bit jump address $\text{Address}_{26\text{ bits}}$
 - 00
- Need an extra control signal decoded from opcode

258

3/2/2016 Chapter 4-The Processor- RST@NTHU

Datapath With Jumps Added



259

3/2/2016 Chapter 4-The Processor- RST@NTHU



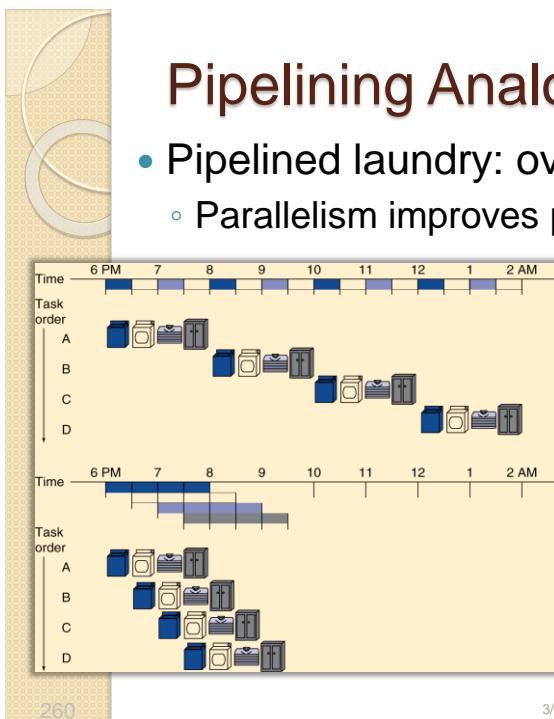
Performance Issues

- Longest delay determines clock period
 - Critical path: load instruction
 - Instruction memory → register file → ALU → data memory → register file
- Not feasible to vary period for different instructions
- Violates design principle
 - Making the common case fast
- We will improve performance by **pipelining**

259

3/2/2016 Chapter 4-The Processor- RST@NTHU

§4.5 An Overview of Pipelining

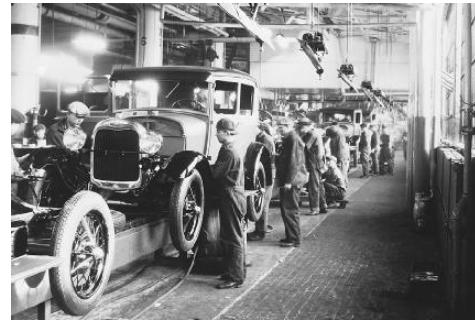
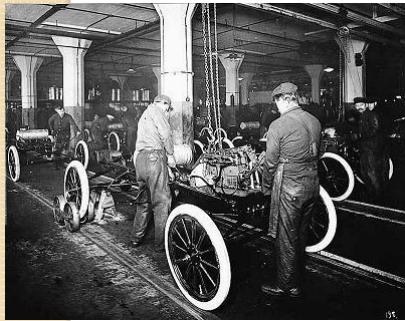


260

3/2/2016 Chapter 4-The Processor- RST@NTHU

- Four loads:
 - $4 * 4 = 16$ hours
- Pipeline non-stop:
 - Speedup
 $= (4 * 4) / 7 = 2.3$
 - $4n / (1n + 3) \approx 4$
= number of stages

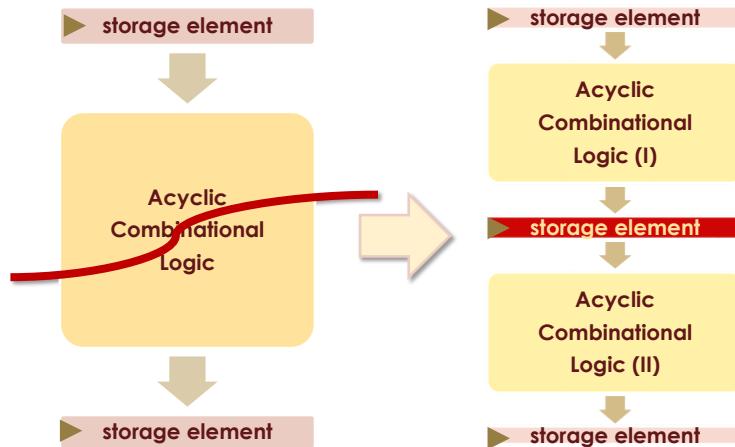
Ford's Assembly Line



3/2/2016 Chapter 4-The Processor- RST@NTHU

Multicycle Implementation

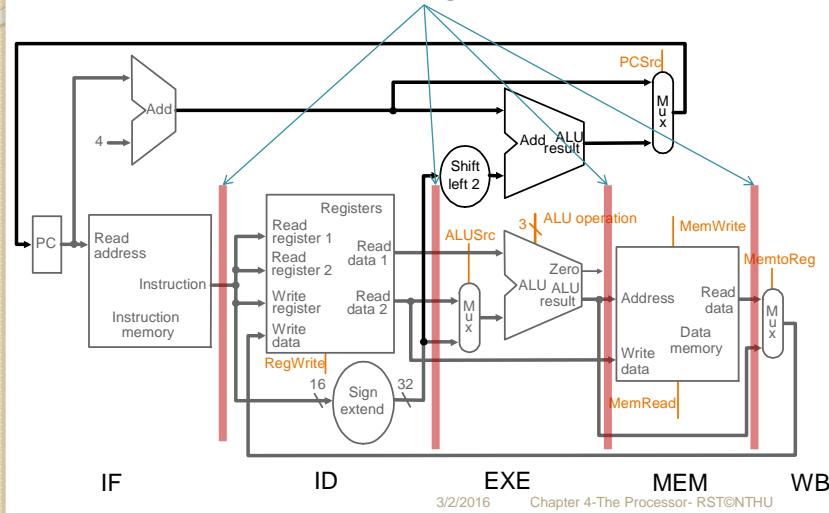
- Balance the work amount of each partition



3/2/2016 Chapter 4-The Processor- RST@NTHU

Partition Single-Cycle Datapath

- Add intermediate registers



MIPS Pipeline

Five stages, one step per stage

- IF:** Instruction fetch from memory
- ID:** Instruction decode & register read
- EX:** Execute operation or calculate address
- MEM:** Access memory operand
- WB:** Write result back to register

Micro-Instructions

Step name	Action for R-type instructions	Action for memory-reference instructions		Action for branches	Action for jumps
IF		IR = Memory[PC] PC = PC + 4			
ID		A = Reg [IR[25-21]] B = Reg [IR[20-16]]			
EXE	ALUOut = A op B	ALUOut = A + sign-extend (IR[15-0])	if (A == B) then ALUOut = PC + (sign-extend (IR[15-0]) << 2) PC = ALUOut	PC = PC [31-28] (IR[25-0]<<2)	
MEM		Load: MDR = Memory[ALUOut] or Store: Memory [ALUOut] = B			
WB	Reg [IR[15-11]] = ALUOut	Load: Reg[IR[20-16]] = MDR			

3/2/2016 Chapter 4-The Processor RST©NTHU

Pipeline Performance

- Assume time for stages is
 - 100ps for register read or write
 - 200ps for other stages
- Compare pipelined datapath with single-cycle datapath

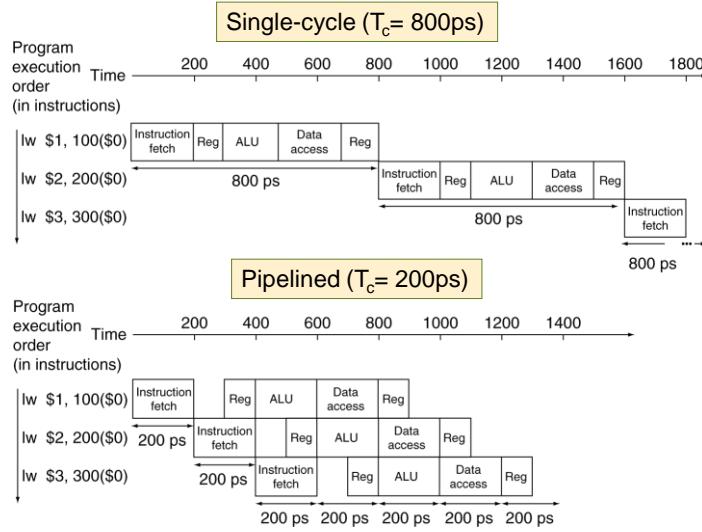
Instr	Instr fetch	Register read	ALU op	Memory access	Register write	Total time
lw	200ps	100 ps	200ps	200ps	100 ps	800ps
sw	200ps	100 ps	200ps	200ps		700ps
R-format	200ps	100 ps	200ps		100 ps	600ps
beq	200ps	100 ps	200ps			500ps

263

3/2/2016 Chapter 4-The Processor RST©NTHU



Pipeline Performance



264

3/2/2016 Chapter 4-The Processor- RST@NTHU



Pipeline Speedup

- If all stages are balanced
 - i.e., all take the same time
 - Time between instructions_{pipelined} = Time between instructions_{nonpipelined} / Number of stages
- If not balanced, speedup is less
- Speedup due to increased throughput
 - Latency (time for each instruction) does not decrease

3/2/2016 Chapter 4-The Processor- RST@NTHU

Pipelining and ISA Design

- MIPS ISA designed for pipelining
 - All instructions are 32-bits
 - Easier to fetch and decode in one cycle
 - c.f. x86: 1- to 17-byte instructions
 - Few and regular instruction formats
 - Can decode and read registers in one step
 - Load/store addressing
 - Can calculate address in 3rd stage, access memory in 4th stage (IF, ID, EXE, MEM, WB)
 - Alignment of memory operands
 - Memory access takes only one cycle

3/2/2016 Chapter 4-The Processor- RST@NTHU

Hazards

- Situations that prevent starting next instruction in next cycle
- **Structure** hazards
 - A required resource is busy
- **Data** hazard
 - Need to wait for previous instruction to complete its data read/write
- **Control** hazard
 - Deciding on control action depends on previous instruction

265

3/2/2016 Chapter 4-The Processor- RST@NTHU



Structure Hazards

- Conflict for use of a resource
- In MIPS pipeline with a single memory
 - Load/store requires data access
 - Instruction fetch would have to *stall* for that cycle
 - Would cause a pipeline “bubble”
- Hence, pipelined datapaths require **separate** instruction/data memories
 - Or separate instruction/data caches

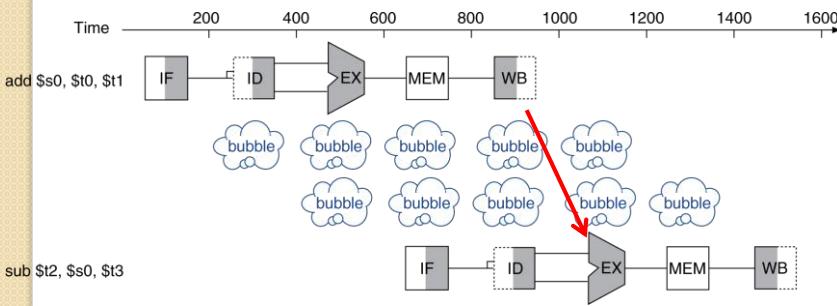
265

3/2/2016 Chapter 4-The Processor- RST@NTHU



Data Hazards

- An instruction depends on completion of data access by a previous instruction
 - add \$s0, \$t0, \$t1
 - sub \$t2, \$s0, \$t3

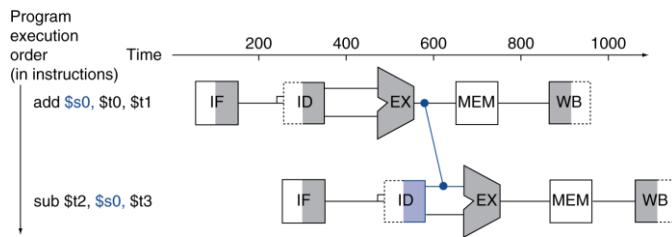


266

3/2/2016 Chapter 4-The Processor- RST@NTHU

Forwarding (aka Bypassing)

- Use result when it is computed
 - Don't wait for it to be stored in a register
 - Requires extra connections in the datapath

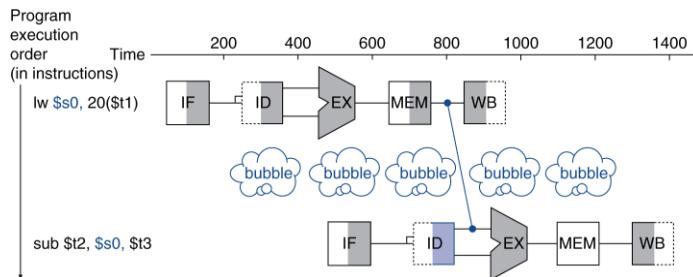


266

3/2/2016 Chapter 4-The Processor- RST@NTHU

Load-Use Data Hazard

- Cannot always avoid stalls by forwarding
 - If value not computed when needed
 - Cannot forward in time!

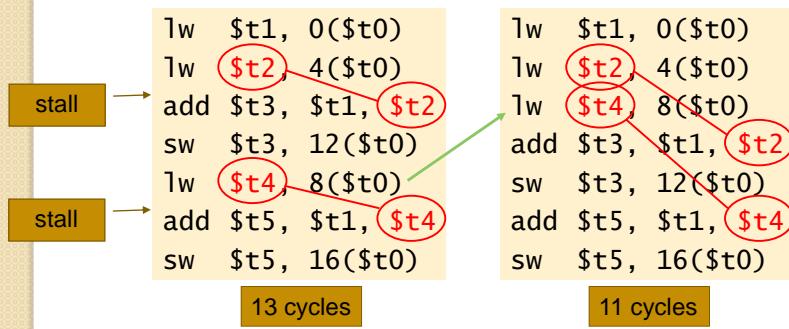


268

3/2/2016 Chapter 4-The Processor- RST@NTHU

Code Scheduling to Avoid Stalls

- Reorder code to avoid use of load result in the next instruction
- C code for $A = B + E; C = B + F;$



268

3/2/2016 Chapter 4-The Processor- RST@NTHU

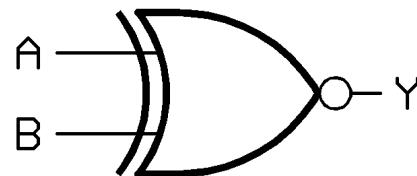
Control Hazards

- Branch determines flow of control
 - Fetching next instruction depends on branch outcome
 - Pipeline can't always fetch correct instruction
 - Still working on ID stage of branch
- In MIPS pipeline
 - Need to compare registers and compute target early in the pipeline
 - Add hardware to do it in ID stage

269

3/2/2016 Chapter 4-The Processor- RST@NTHU

Check EQ using XNOR

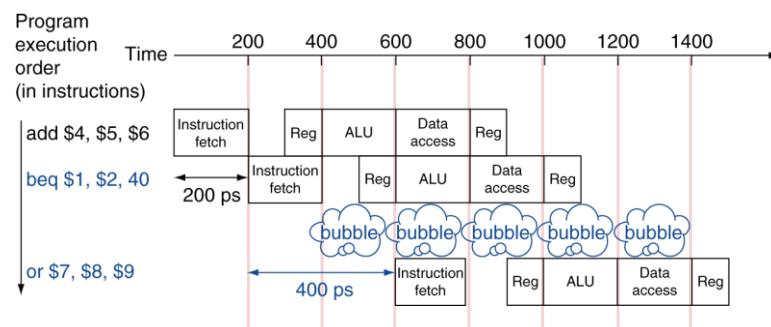


A	B	Y
0	0	1
0	1	0
1	0	0
1	1	1

3/2/2016 Chapter 4-The Processor- RST@NTU

Stall on Branch

- Wait until branch outcome determined before fetching next instruction



270

3/2/2016 Chapter 4-The Processor- RST@NTU



Branch Prediction

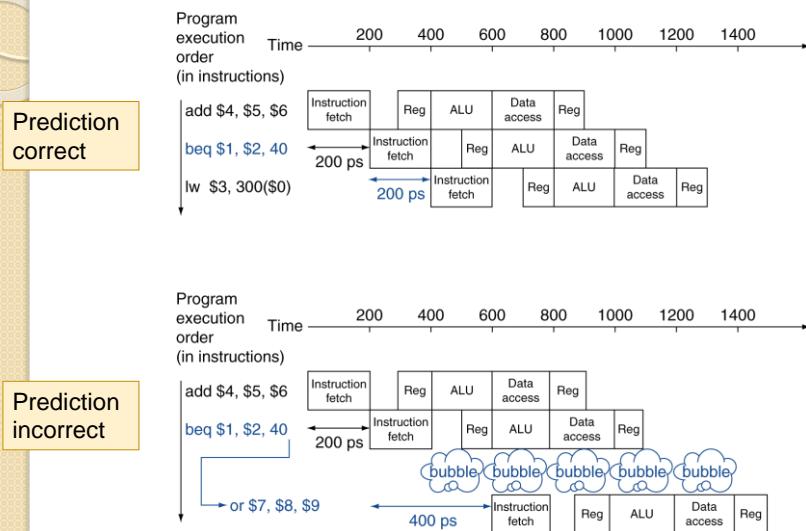
- Longer pipelines can't readily determine branch outcome early
 - Stall penalty becomes unacceptable
- Predict outcome of branch
 - Only stall if prediction is wrong
- In MIPS pipeline
 - Can predict branches not taken
 - Fetch instruction after branch, with no delay

271

3/2/2016 Chapter 4-The Processor- RST@NTHU



MIPS with Predict Not Taken



272

3/2/2016 Chapter 4-The Processor- RST@NTHU



More-Realistic Branch Prediction

- Static branch prediction
 - Based on typical branch behavior
 - Example: loop and if-statement branches
 - Predict backward branches taken
 - Predict forward branches not taken
- Dynamic branch prediction
 - Hardware measures actual branch behavior
 - e.g., record recent history of each branch
 - Assume future behavior will continue the trend
 - When wrong, stall while re-fetching, and update history



3/2/2016 Chapter 4-The Processor- RST@NTHU



Pipeline Summary

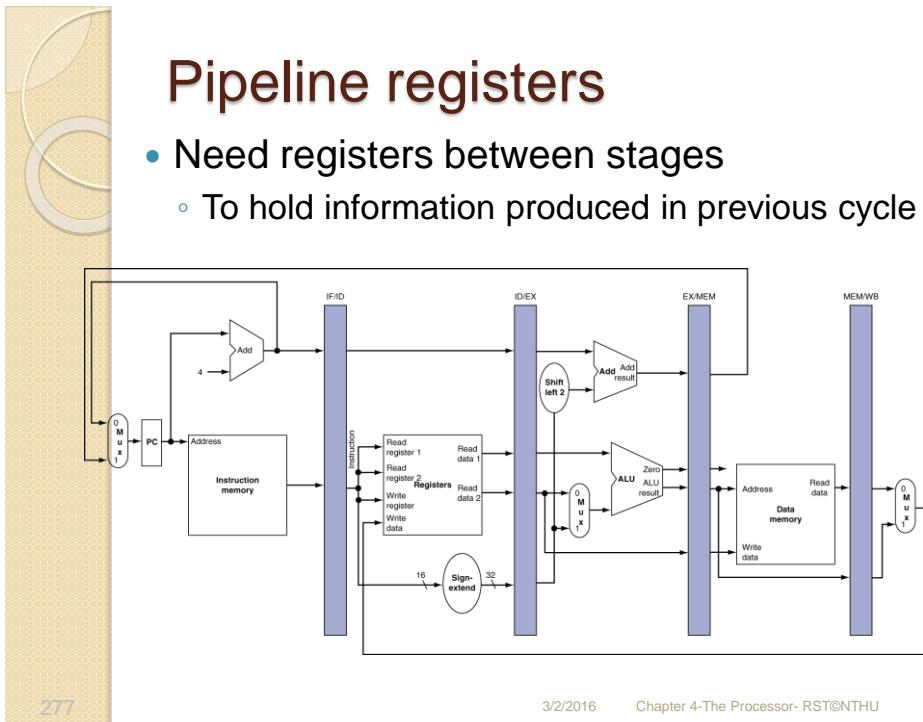
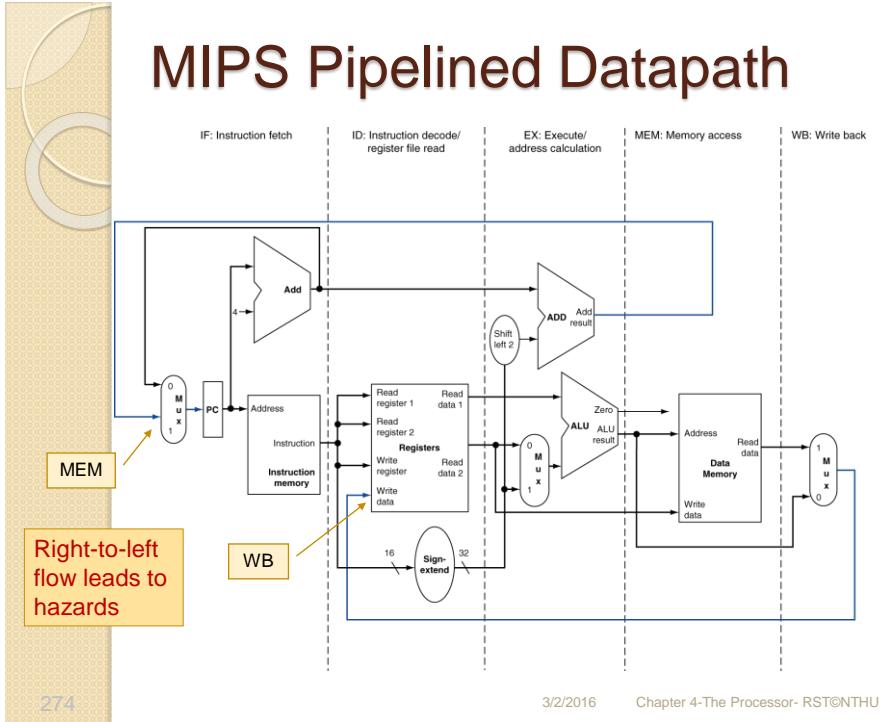
The BIG Picture

- Pipelining improves performance by increasing instruction throughput
 - Executes multiple instructions in parallel
 - Each instruction has the same latency
- Subject to hazards
 - Structure, data, control
- Instruction set design affects complexity of pipeline implementation

273

3/2/2016 Chapter 4-The Processor- RST@NTHU

§4.6 Pipelined Datapath and Control

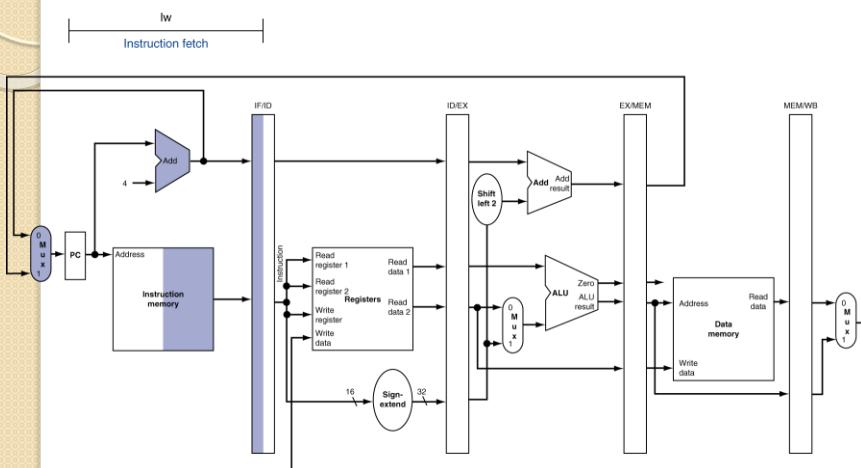


Pipeline Operation

- Cycle-by-cycle flow of instructions through the pipelined datapath
 - “Single-clock-cycle” pipeline diagram
 - Shows pipeline usage in a single cycle
 - Highlight resources used
 - c.f. “multi-clock-cycle” diagram
 - Graph of operation over time
- We'll look at “single-clock-cycle” diagrams for load & store

3/2/2016 Chapter 4-The Processor- RST@NTHU

IF for Load, Store, ...

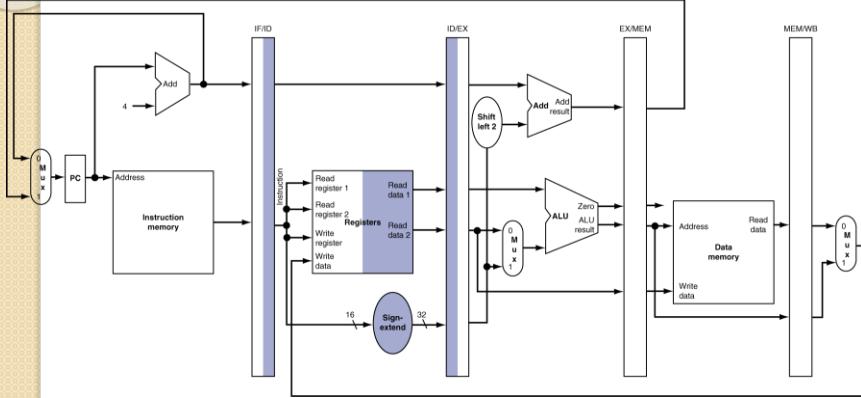


279

3/2/2016 Chapter 4-The Processor- RST@NTHU

ID for Load, Store, ...

lw
Instruction decode

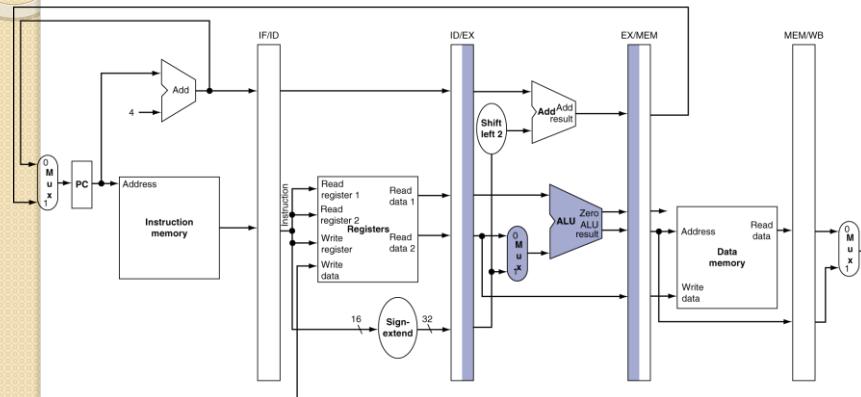


279

3/2/2016 Chapter 4-The Processor- RST@NTHU

EX for Load

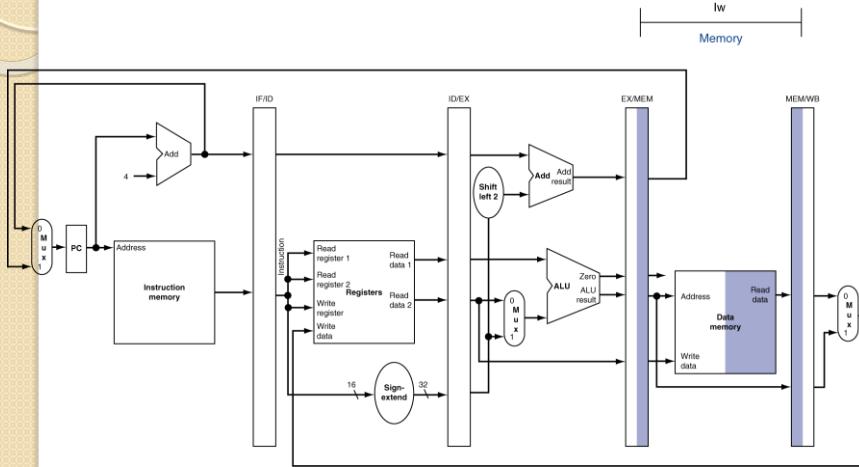
lw
Execution



280

3/2/2016 Chapter 4-The Processor- RST@NTHU

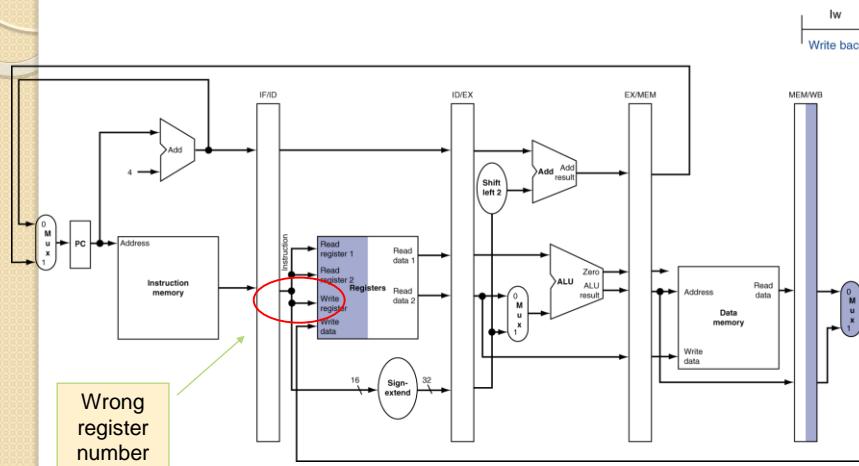
MEM for Load



281

3/2/2016 Chapter 4-The Processor- RST@NTHU

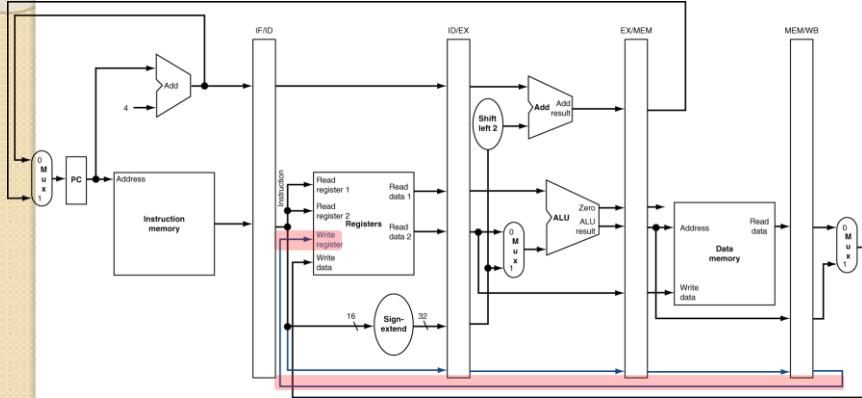
WB for Load



281

3/2/2016 Chapter 4-The Processor- RST@NTHU

Corrected Datapath for Load

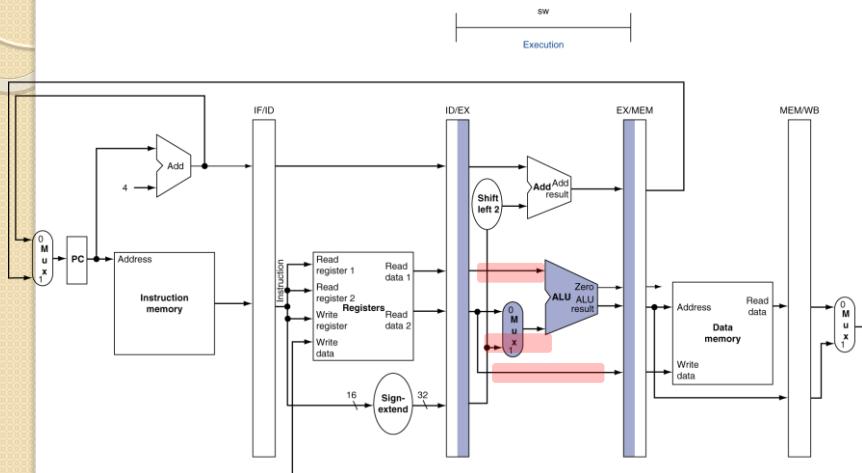


284

3/2/2016 Chapter 4-The Processor- RST@NTHU

EX for Store

sw \$t0, 16(\$s1)

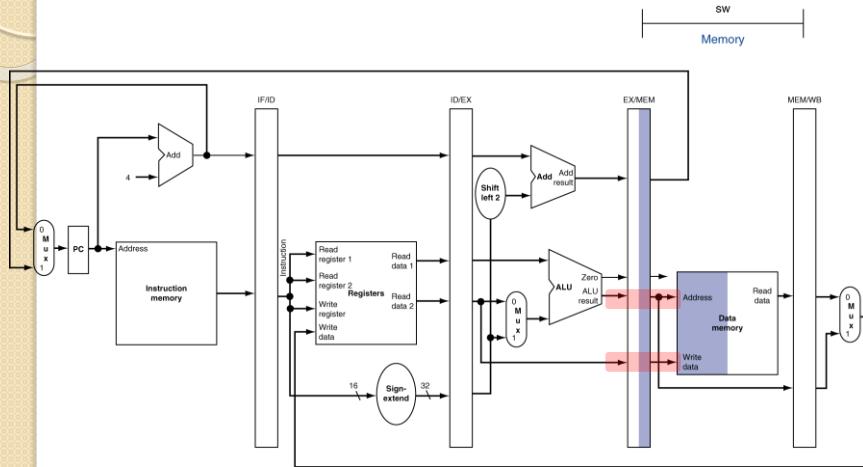


282

3/2/2016 Chapter 4-The Processor- RST@NTHU

MEM for Store

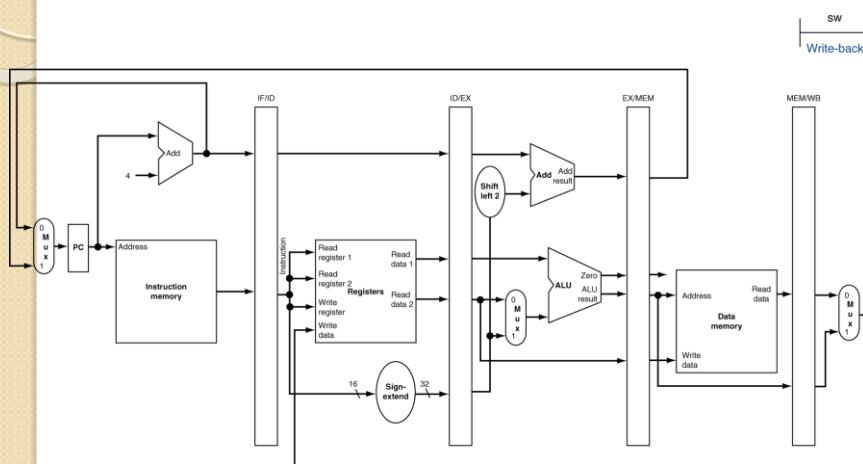
sw \$t0, 16(\$s1)



283

3/2/2016 Chapter 4-The Processor- RST@NTHU

WB for Store



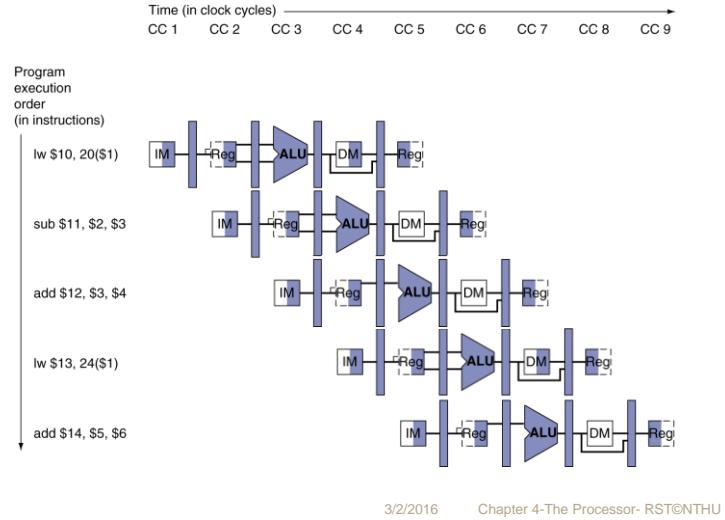
283

3/2/2016 Chapter 4-The Processor- RST@NTHU



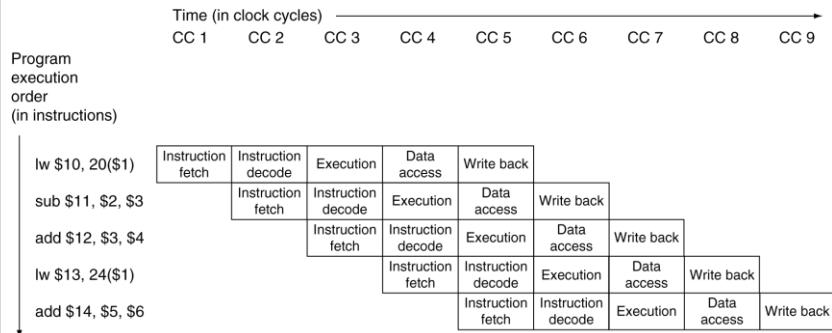
Multi-Cycle Pipeline Diagram

- Form showing **resource** usage



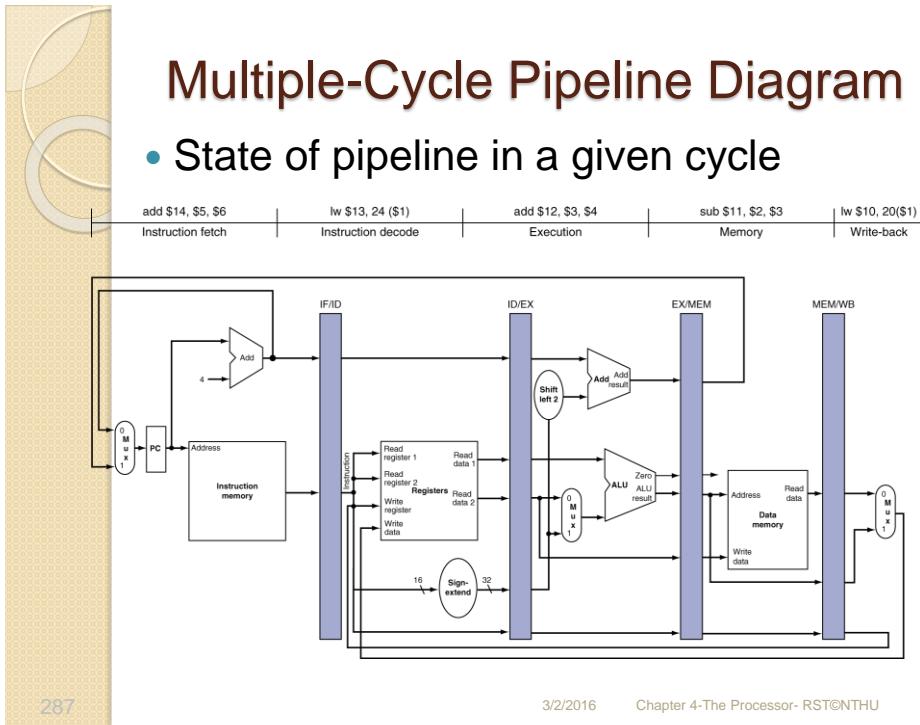
Multi-Cycle Pipeline Diagram

- Traditional** form



Multiple-Cycle Pipeline Diagram

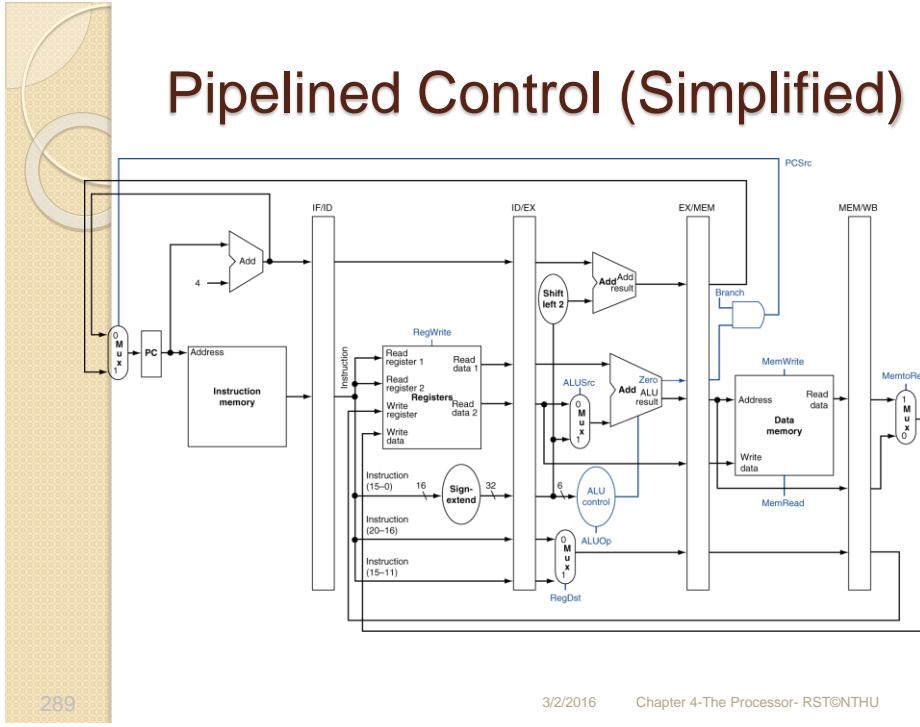
- State of pipeline in a given cycle



287

3/2/2016 Chapter 4-The Processor- RST@NTHU

Pipelined Control (Simplified)

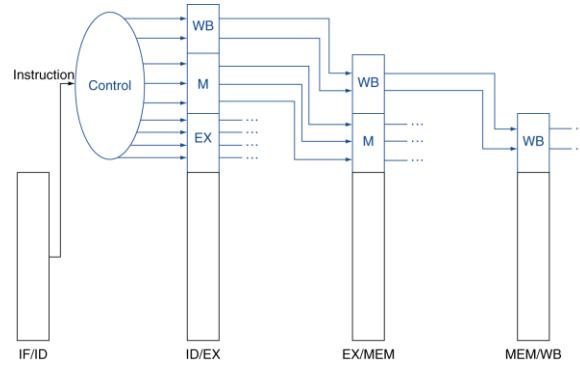


289

3/2/2016 Chapter 4-The Processor- RST@NTHU

Pipelined Control

- Control signals derived from instruction
 - As in single-cycle implementation

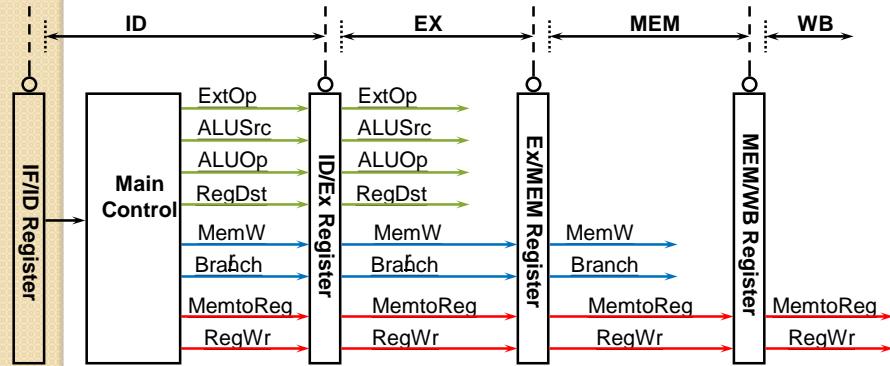


291

3/2/2016 Chapter 4-The Processor- RST@NTHU

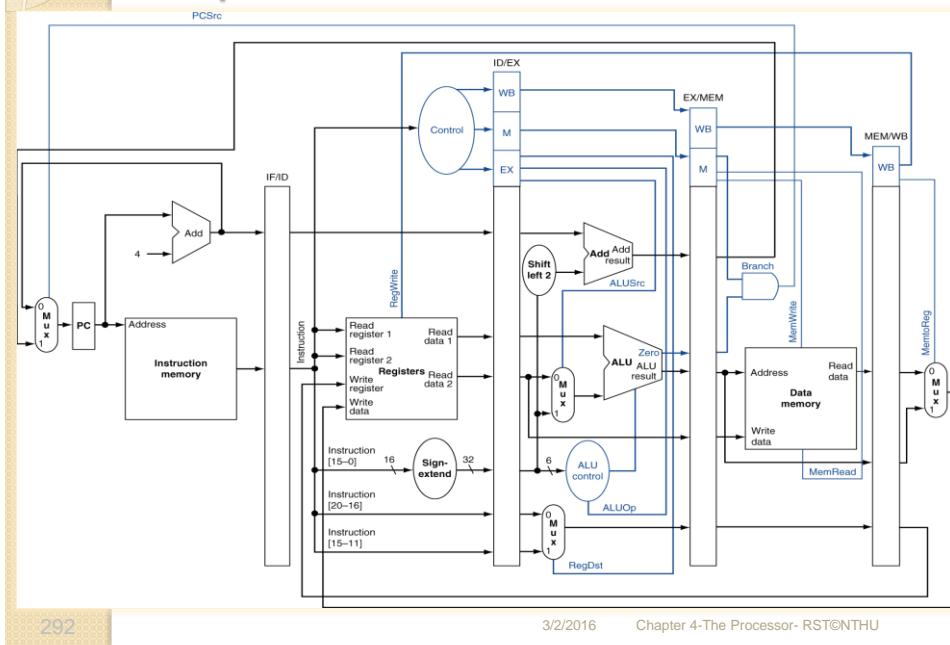
Pipelined Control (cont.)

- Signals for EX (ExtOp, ALUSrc, ...) are used 1 cycle later
- Signals for MEM (MemWr, Branch) are used 2 cycles later
- Signals for WB (MemtoReg, RegWr) are used 3 cycles later



3/2/2016 Chapter 4-The Processor- RST@NTHU

Pipelined Control



292

3/2/2016 Chapter 4-The Processor- RST@NTHU

Data Hazards in ALU Instructions

§4.7 Data Hazards Forwarding vs. Stalling

- Consider this sequence:


```
sub $2, $1,$3
and $12,$2,$5
or $13,$6,$2
add $14,$2,$2
sw $15,100($2)
```
- We can resolve hazards with forwarding
 - How do we detect when to forward?

291

3/2/2016 Chapter 4-The Processor- RST@NTHU

Data Dependency & Availability

- R-type

op	rs	rt	rd	shamt	funct
----	----	----	----	-------	-------

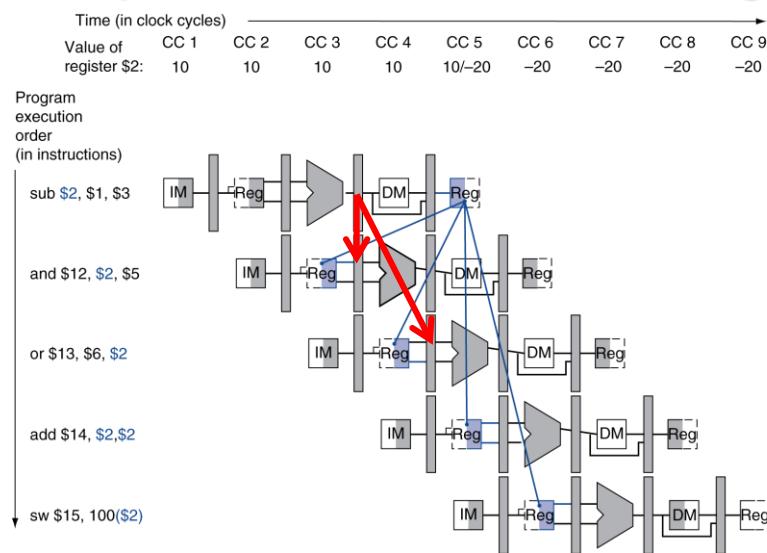
 - Inputs: rs, rt @ ID/EX
 - Output: rd @ EX/MEM
- I-type

op	rs	rt	constant or address		
----	----	----	---------------------	--	--

 - lw
 - Inputs: rs @ ID/EX
 - Output: rt @ MEM/WB
 - SW:
 - Inputs: rs @ ID/EX, rt @ EX/MEM
 - branch
 - Inputs: rs, rt @ ID
 - Output: PC @ ID (move from EX)

3/2/2016 Chapter 4-The Processor- RST@NTHU

Dependencies & Forwarding



3/2/2016 Chapter 4-The Processor- RST@NTHU

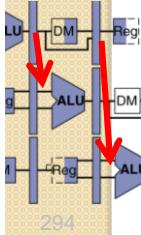


Detecting the Need to Forward

- Pass register numbers along pipeline
 - e.g., ID/EX.RegisterRs = register number for Rs sitting in ID/EX pipeline register
- ALU operand register numbers in EX stage are given by
 - ID/EX.RegisterRs, ID/EX.RegisterRt
- Data hazards when

- | | |
|--|--|
| 1a. EX/MEM.RegisterRd = ID/EX.RegisterRs
1b. EX/MEM.RegisterRd = ID/EX.RegisterRt
2a. MEM/WB.RegisterRd = ID/EX.RegisterRs
2b. MEM/WB.RegisterRd = ID/EX.RegisterRt | Fwd from EX/MEM pipeline reg

Fwd from MEM/WB pipeline reg |
|--|--|



3/2/2016 Chapter 4-The Processor- RST@NTHU

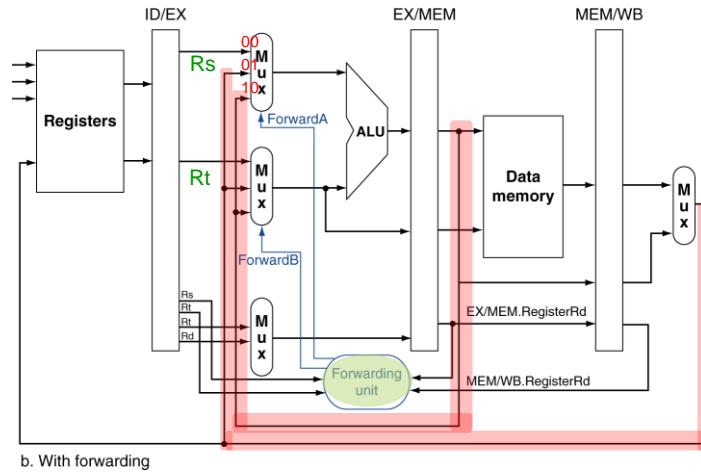


Detecting the Need to Forward

- But only if forwarding instruction will write to a register!
 - EX/MEM.RegWrite, MEM/WB.RegWrite
- And only if Rd for that instruction is not \$zero
 - EX/MEM.RegisterRd ≠ 0, MEM/WB.RegisterRd ≠ 0

3/2/2016 Chapter 4-The Processor- RST@NTHU

Forwarding Paths



b. With forwarding

297

3/2/2016 Chapter 4-The Processor- RST@NTHU

Forwarding Conditions

- EX hazard (EX/MEM)
 - if (EX/MEM.RegWrite and (EX/MEM.RegisterRd ≠ 0)
and (EX/MEM.RegisterRd = ID/EX.RegisterRs))
ForwardA = 10
 - if (EX/MEM.RegWrite and (EX/MEM.RegisterRd ≠ 0)
and (EX/MEM.RegisterRd = ID/EX.RegisterRt))
ForwardB = 10
- MEM hazard (MEM/WB)
 - if (MEM/WB.RegWrite and (MEM/WB.RegisterRd ≠ 0)
and (MEM/WB.RegisterRd = ID/EX.RegisterRs))
ForwardA = 01
 - if (MEM/WB.RegWrite and (MEM/WB.RegisterRd ≠ 0)
and (MEM/WB.RegisterRd = ID/EX.RegisterRt))
ForwardB = 01

296,298

3/2/2016 Chapter 4-The Processor- RST@NTHU



Double Data Hazard

- Consider the sequence:
$$\begin{array}{l} \text{add \$1,\$1,\$2} \\ \text{add \$1,\$1,\$3} \\ \text{add \$1,\$1,\$4} \end{array}$$
 } } *
- Both hazards occur
 - Want to use the most recent
- Revise MEM hazard condition
 - fwd only if EX hazard condition isn't true

298

3/2/2016 Chapter 4-The Processor- RST@NTHU



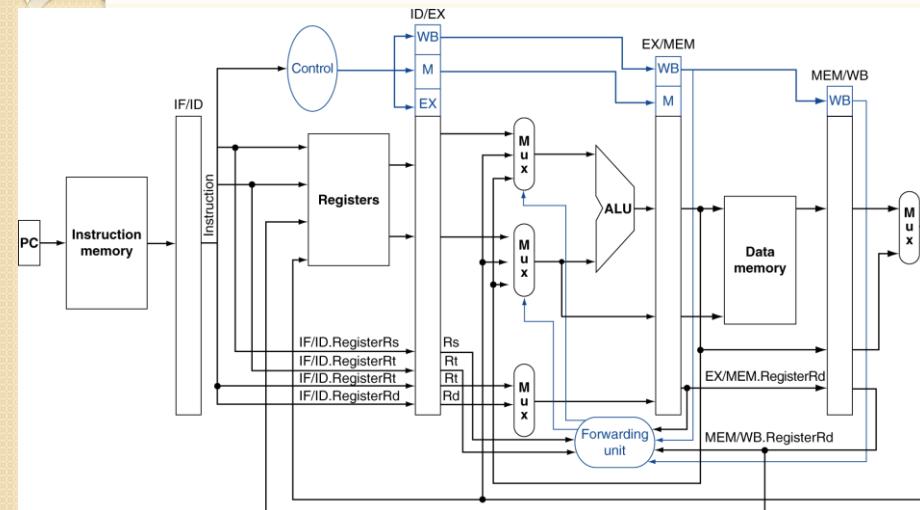
Revised Forwarding Condition

- MEM hazard
 - if (MEM/WB.RegWrite and (MEM/WB.RegisterRd ≠ 0)
and not (EX/MEM.RegWrite and (EX/MEM.RegisterRd ≠ 0)
and (EX/MEM.RegisterRd = ID/EX.RegisterRs))
and (MEM/WB.RegisterRd = ID/EX.RegisterRs))
ForwardA = 01
 - if (MEM/WB.RegWrite and (MEM/WB.RegisterRd ≠ 0)
and not (EX/MEM.RegWrite and (EX/MEM.RegisterRd ≠ 0)
and (EX/MEM.RegisterRd = ID/EX.RegisterRt))
and (MEM/WB.RegisterRd = ID/EX.RegisterRt))
ForwardB = 01

299

3/2/2016 Chapter 4-The Processor- RST@NTHU

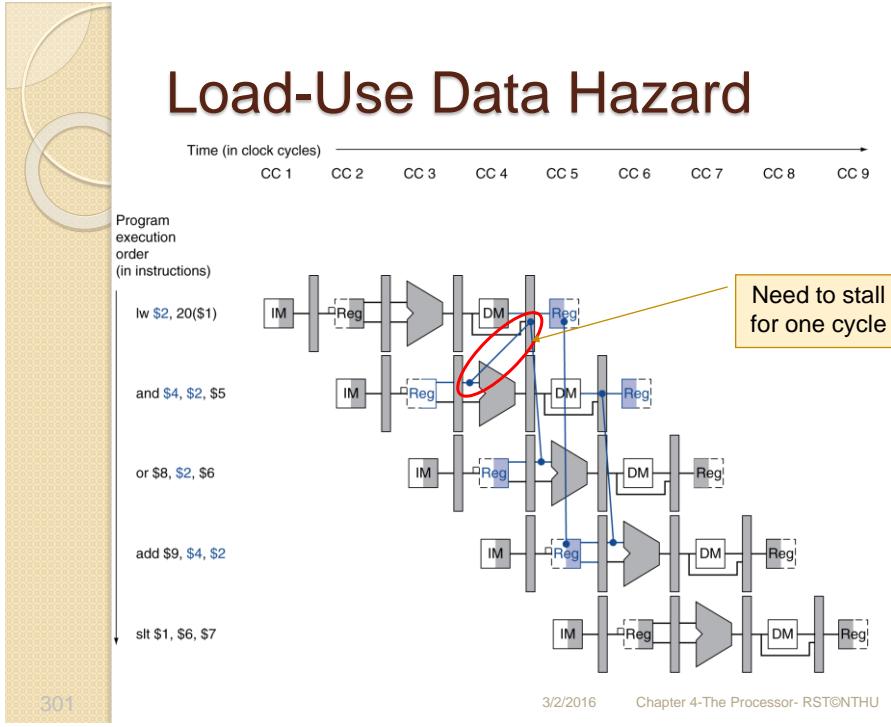
Datapath with Forwarding



299

3/2/2016 Chapter 4-The Processor- RST@NTHU

Load-Use Data Hazard



301

3/2/2016 Chapter 4-The Processor- RST@NTHU



Load-Use Hazard Detection

- Check when using instruction is decoded in ID stage
- ALU operand register numbers in ID stage are given by
 - IF/ID.RegisterRs, IF/ID.RegisterRt
- Load-use hazard when
 - ID/EX.MemRead and
 $((ID/EX.RegisterRt = IF/ID.RegisterRs)$
or
 $(ID/EX.RegisterRt = IF/ID.RegisterRt))$
- If detected, stall and insert bubble

3/2/2016 Chapter 4-The Processor- RST@NTHU



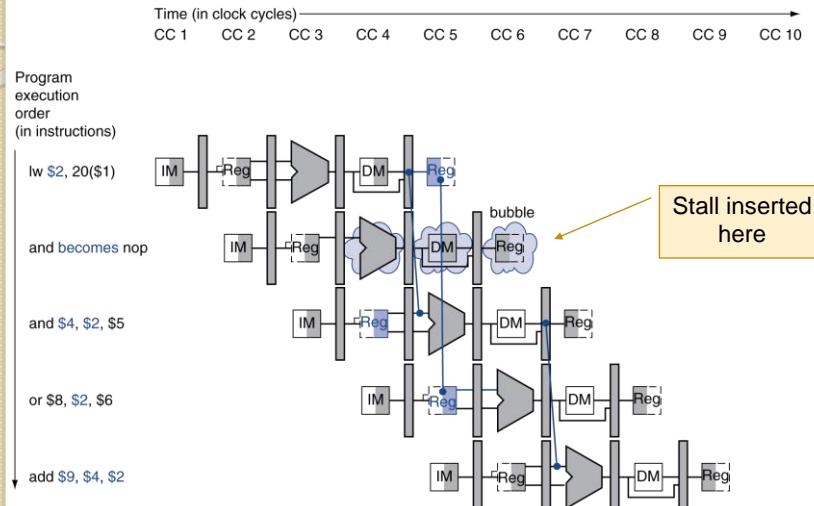
How to Stall the Pipeline

- Force control values in ID/EX register to 0
 - EX, MEM and WB do **nop** (no-operation)
- Prevent update of PC and IF/ID register
 - Using instruction is decoded again
 - Following instruction is fetched again
 - 1-cycle stall allows MEM to read data for 1w
 - Can subsequently forward to EX stage

3/2/2016 Chapter 4-The Processor- RST@NTHU



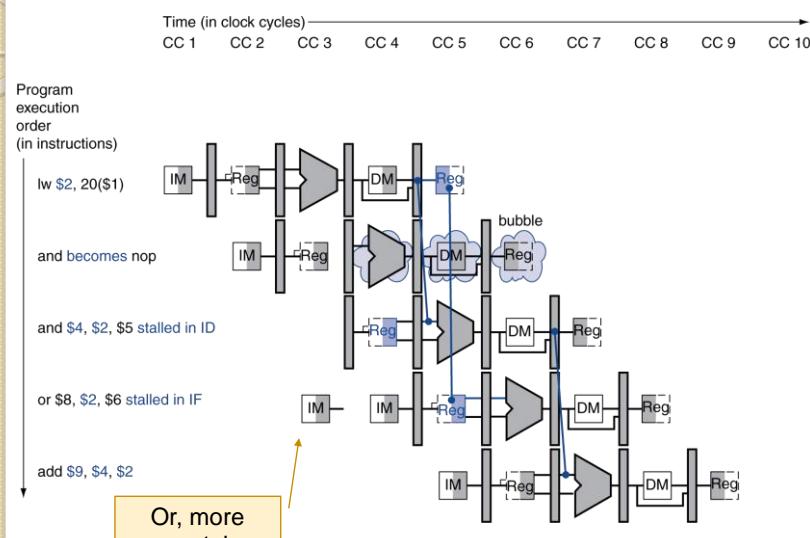
Stall/Bubble in the Pipeline



3/2/2016 Chapter 4-The Processor- RST@NTHU

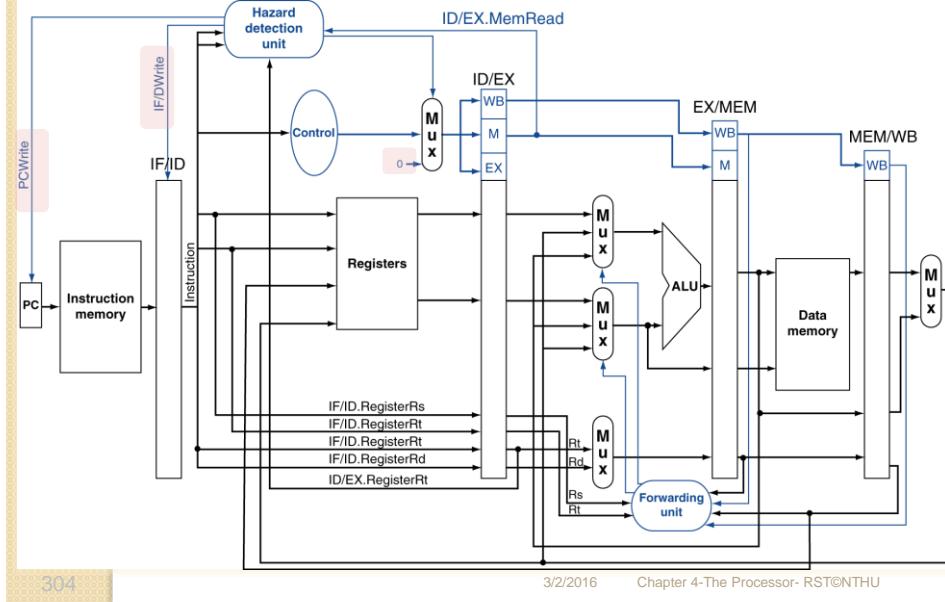


Stall/Bubble in the Pipeline



3/2/2016 Chapter 4-The Processor- RST@NTHU

Datapath with Hazard Detection



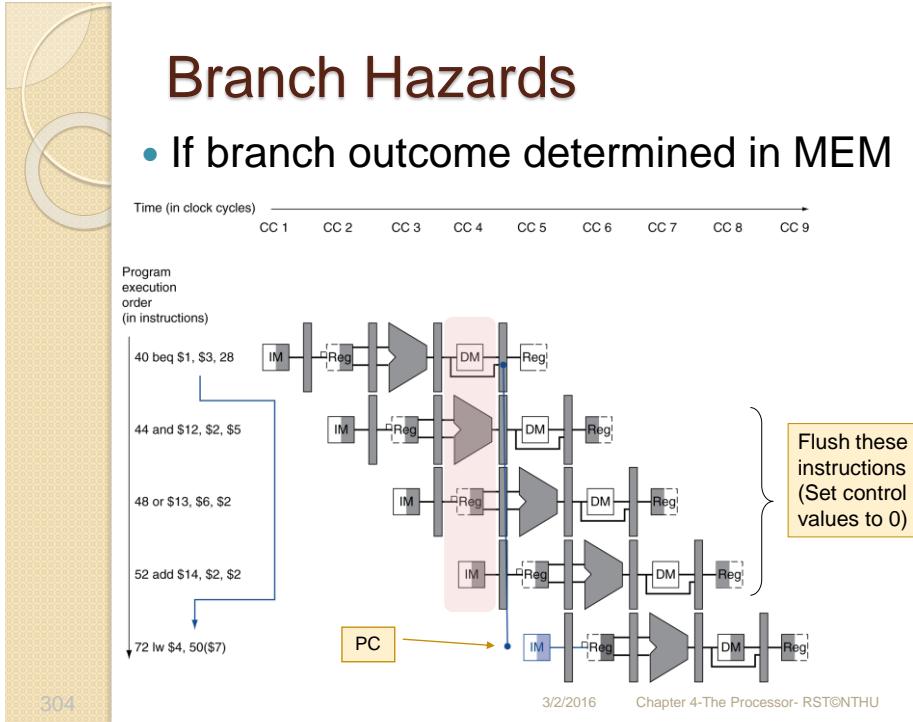
Stalls and Performance

The BIG Picture

- Stalls reduce performance
 - But are required to get correct results
- Compiler can arrange code to avoid hazards and stalls
 - Requires knowledge of the pipeline structure

Branch Hazards

- If branch outcome determined in MEM



304

Reducing Branch Delay

- Move hardware to determine outcome to ID stage
 - Target address adder
 - Register comparator
- Example: branch taken

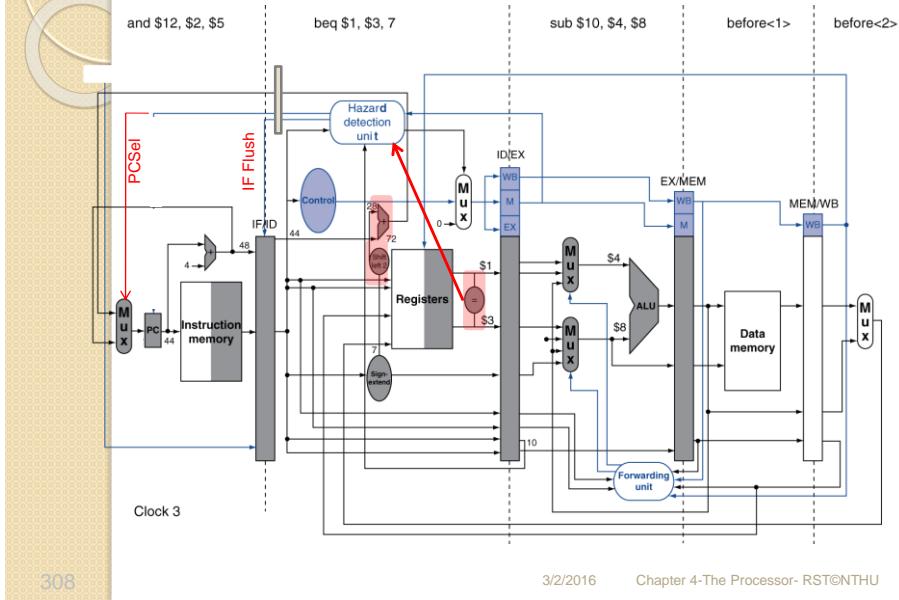
```

36: sub $10, $4, $8
40: beq $1, $3, 7 # (40+4)+7*4=72
44: and $12, $2, $5
48: or $13, $2, $6
52: add $14, $4, $2
56: slt $15, $6, $7
      ...
72: lw $4, 50($7)
  
```



Move
comparator
to ID

Example: Branch Taken 1/2

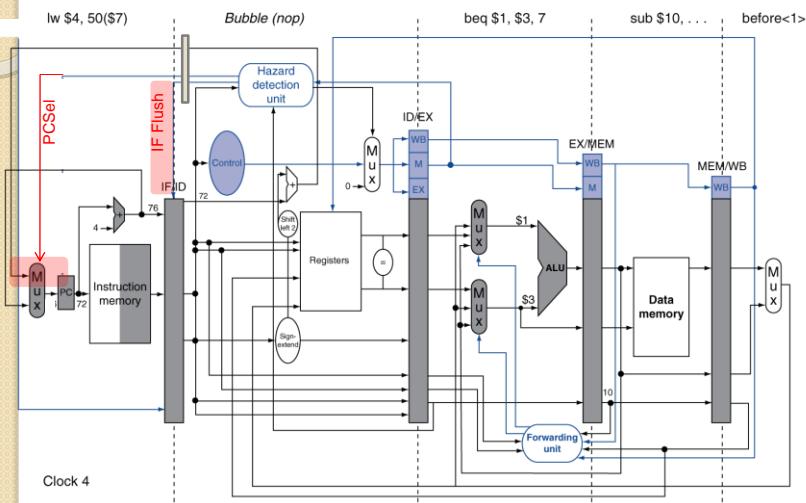


308

3/2/2016 Chapter 4-The Processor- RST@NTHU



Example: Branch Taken 2/2

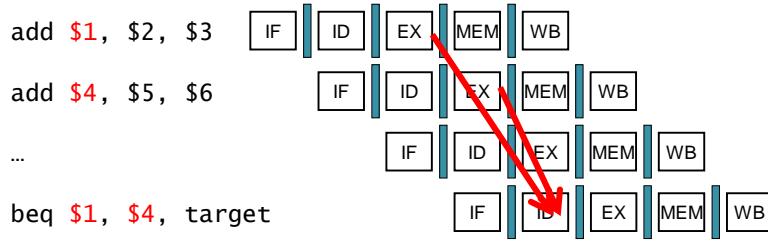


308

3/2/2016 Chapter 4-The Processor- RST@NTHU

Data Hazards for Branches

- If a comparison register is a destination of 2nd or 3rd preceding ALU instruction

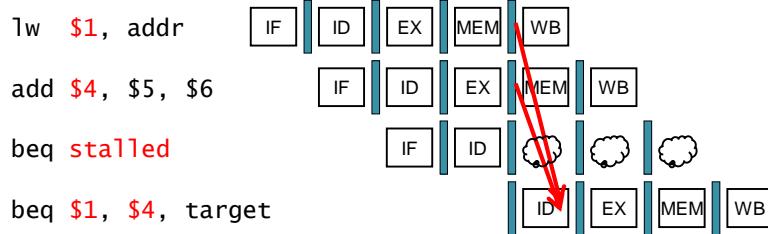


- Can resolve using forwarding

3/2/2016 Chapter 4-The Processor- RST@NTHU

Data Hazards for Branches

- If a comparison register is a destination of preceding ALU instruction or 2nd preceding load instruction
 - Need 1 stall cycle



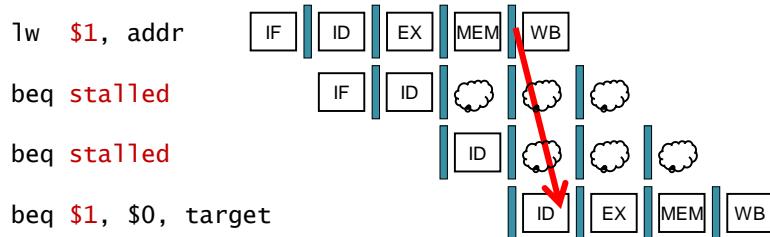
(what if branch address is calculated using ALU?)

3/2/2016

Chapter 4-The Processor- RST@NTHU

Data Hazards for Branches

- If a comparison register is a destination of immediately preceding load instruction
 - Need 2 stall cycles



3/2/2016 Chapter 4-The Processor- RST@NTHU

Dynamic Branch Prediction

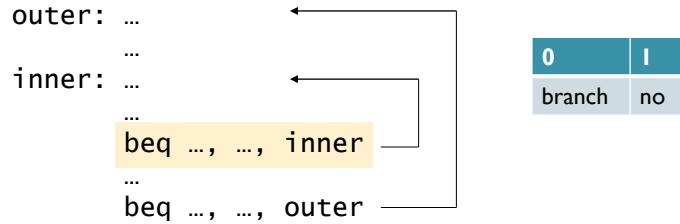
- In deeper and superscalar pipelines, branch penalty is more significant
- Use dynamic prediction
 - Branch prediction buffer (aka branch history table)
 - Indexed by recent branch instruction addresses
 - Stores outcome (taken/not taken)
 - To execute a branch
 - Check table, expect the same outcome
 - Start fetching from fall-through or target
 - If wrong, flush pipeline and flip prediction

309

3/2/2016 Chapter 4-The Processor- RST@NTHU

1-Bit Predictor: Shortcoming

- Inner loop branches mispredicted twice!

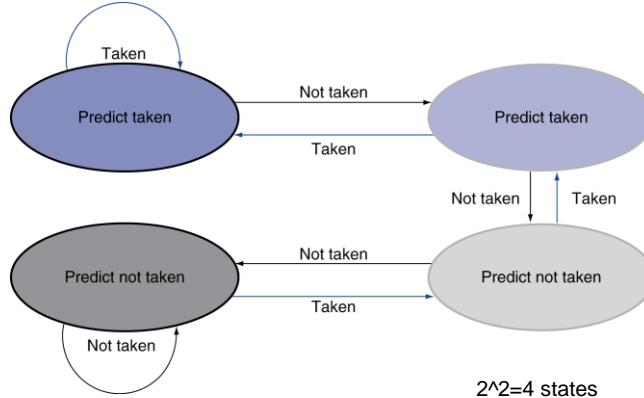


- Mispredict as taken on last iteration of inner loop
- Then mispredict as not taken on first iteration of inner loop next time around

3/2/2016 Chapter 4-The Processor- RST@NTHU

2-Bit Predictor

- Only change prediction on two successive mispredictions



310

3/2/2016 Chapter 4-The Processor- RST@NTHU

Calculating the Branch Target

- Even with predictor, still need to calculate the target address
 - 1-cycle penalty for a taken branch
- Branch target buffer
 - Cache of target addresses
 - Indexed by PC when instruction fetched
 - If hit and instruction is branch predicted taken, can fetch target immediately

3/2/2016 Chapter 4-The Processor- RST@NTHU

§4.9 Exceptions

Exceptions and Interrupts

- “Unexpected” events requiring change in flow of control
 - Different ISAs use the terms differently
- Exception
 - Arises within the CPU
 - e.g., undefined opcode, overflow, syscall, ...
- Interrupt
 - From an external I/O controller
- Dealing with them without sacrificing performance is hard

313

3/2/2016 Chapter 4-The Processor- RST@NTHU



Handling Exceptions

- In MIPS, exceptions managed by a System Control Coprocessor (CP0)
- Save PC of offending (or interrupted) instruction
 - In MIPS: Exception Program Counter (EPC)
- Save indication of the problem
 - In MIPS: Cause register
 - We'll assume 1-bit
 - 0 for undefined opcode, 1 for overflow
- Jump to handler at 8000 0180

314

3/2/2016 Chapter 4-The Processor- RST@NTHU



An Alternate Mechanism

- Vectored Interrupts
 - Handler address determined by the cause
- Example:
 - Undefined opcode: C000 0000
 - Overflow: C000 0020
 -: C000 0040
- Instructions either
 - Deal with the interrupt, or
 - Jump to real handler

3/2/2016 Chapter 4-The Processor- RST@NTHU

Handler Actions

- Read cause, and transfer to relevant handler
- Determine action required
- If restartable
 - Take corrective action
 - use EPC to return to program
- Otherwise
 - Terminate program
 - Report error using EPC, cause, ...

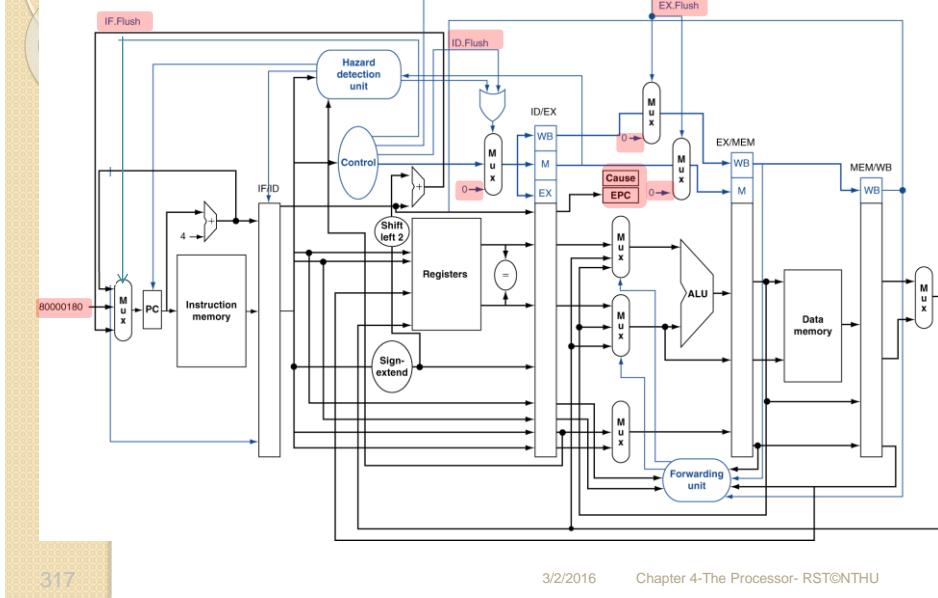
3/2/2016 Chapter 4-The Processor- RST@NTHU

Exceptions in a Pipeline

- Another form of control hazard
- Consider overflow on add in EX stage
add \$1, \$2, \$1
 - Prevent \$1 from being clobbered
 - **Complete previous instructions**
 - Flush add and subsequent instructions
 - Set Cause and EPC register values
 - Transfer control to handler
- Similar to mispredicted branch
 - Use much of the same hardware

3/2/2016 Chapter 4-The Processor- RST@NTHU

Pipeline with Exceptions



317

3/2/2016 Chapter 4-The Processor- RST@NTHU

Exception Properties

- Restartable exceptions
 - Pipeline can flush the instruction
 - Handler executes, then returns to the instruction
 - Refetched and executed from scratch
- PC saved in EPC register
 - Identifies causing instruction
 - Actually PC + 4 is saved
 - Handler must adjust

3/2/2016 Chapter 4-The Processor- RST@NTHU



Exception Example

- Exception on **add** in

```

40    sub   $11, $2, $4
44    and   $12, $2, $5
48    or    $13, $2, $6
4C  add  $1, $2, $1
50    slt   $15, $6, $7
54    lw    $16, 50($7)
...

```

- Handler

```

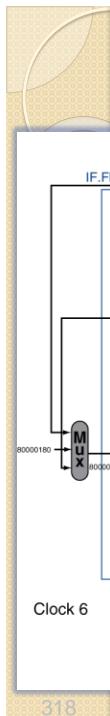
80000180  sw    $26, 1000($0)
80000184  sw    $27, 1004($0)
...

```

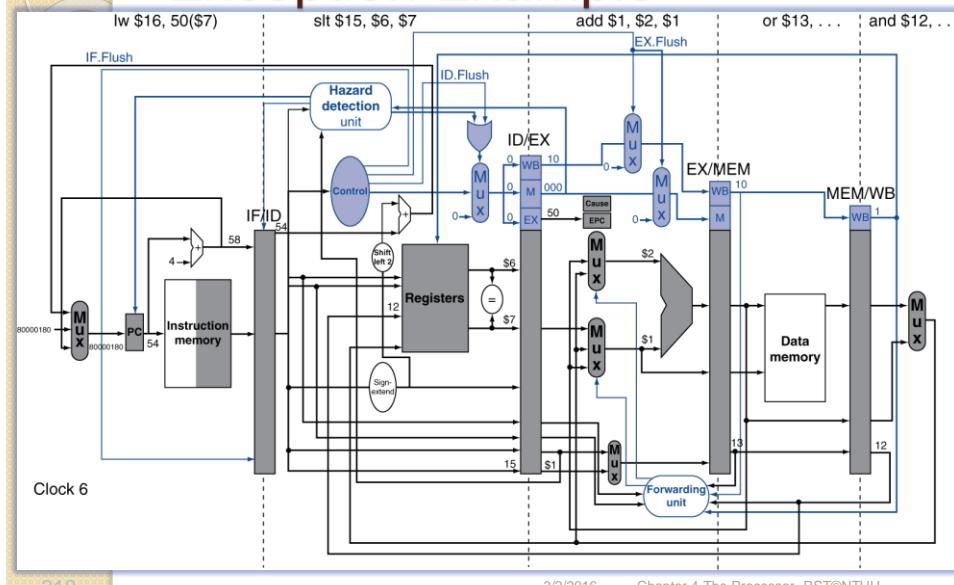
26	k0	for OS
27	k1	kernel

316

3/2/2016 Chapter 4-The Processor- RST@NTHU



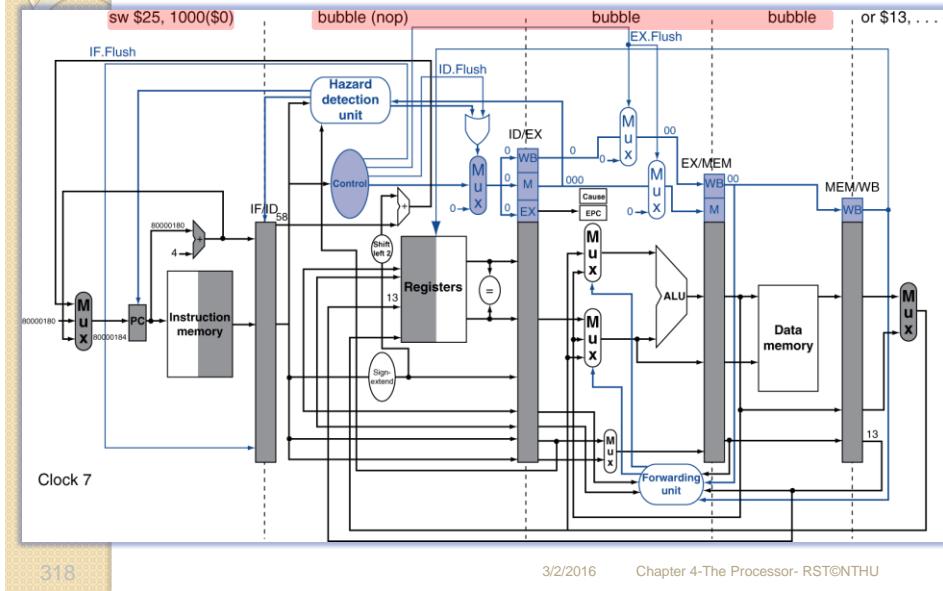
Exception Example



318

3/2/2016 Chapter 4-The Processor- RST@NTHU

Exception Example



3/2/2016 Chapter 4-The Processor- RST@NTHU

Multiple Exceptions

- Pipelining overlaps multiple instructions
 - Could have multiple exceptions at once
- Simple approach: deal with exception of the **earliest** instruction
 - Flush subsequent instructions
 - “Precise” exceptions
- In complex pipelines
 - Multiple instructions issued per cycle
 - Out-of-order completion
 - Maintaining precise exceptions is difficult!

3/2/2016 Chapter 4-The Processor- RST@NTHU



Imprecise Exceptions

- Just stop pipeline and save state
 - Including exception cause(s)
- Let the handler work out
 - Which instruction(s) had exceptions
 - Which to complete or flush
 - May require “manual” completion
- Simplifies hardware, but more complex handler software
- Not feasible for complex multiple-issue out-of-order pipelines

3/2/2016 Chapter 4-The Processor- RST©NTHU



Instruction-Level Parallelism (ILP)

- Pipelining: executing multiple instructions in parallel
- To increase ILP
 - Deeper pipeline
 - Less work per stage \Rightarrow shorter clock cycle
 - Multiple issue
 - Replicate pipeline stages \Rightarrow multiple pipelines
 - Start multiple instructions per clock cycle
 - CPI < 1, so use Instructions Per Cycle (IPC)
 - E.g., 4GHz 4-way multiple-issue
 - 16 BIPS, peak CPI = 0.25, peak IPC = 4
 - But dependencies reduce this in practice

§4.10 Parallelism and Advanced Instruction Level Parallelism

3/2/2016 Chapter 4-The Processor- RST©NTHU



Multiple Issue

- **Static** multiple issue
 - Compiler groups instructions to be issued together
 - Packages them into “issue slots”
 - Compiler detects and avoids hazards
- **Dynamic** multiple issue
 - CPU examines instruction stream and chooses instructions to issue each cycle
 - Compiler can help by reordering instructions
 - CPU resolves hazards using advanced techniques at runtime

320

3/2/2016 Chapter 4-The Processor- RST@NTHU



Speculation

- “Guess” what to do with an instruction
 - Start operation as soon as possible
 - Check whether guess was right
 - If so, complete the operation
 - If not, roll-back and do the right thing
- Common to static and dynamic multiple issue
- Examples
 - Speculate on branch outcome
 - Roll back if path taken is different
 - Speculate on load
 - Roll back if location is updated

321

3/2/2016 Chapter 4-The Processor- RST@NTHU

Compiler/Hardware Speculation

- Compiler can reorder instructions
 - e.g., move load before branch
 - Can include “fix-up” instructions to recover from incorrect guess
- Hardware can look ahead for instructions to execute
 - Buffer results until it determines they are actually needed
 - Flush buffers on incorrect speculation

3/2/2016 Chapter 4-The Processor- RST@NTHU

Speculation and Exceptions

- What if exception occurs on a speculatively executed instruction?
 - e.g., speculative load before null-pointer check
- Static speculation
 - Can add ISA support for deferring exceptions
- Dynamic speculation
 - Can buffer exceptions until instruction completion (which may not occur)

3/2/2016 Chapter 4-The Processor- RST@NTHU



Static Multiple Issue

- Compiler groups instructions into “issue packets”
 - Group of instructions that can be issued on a single cycle
 - Determined by pipeline resources required
- Think of an issue packet as a very long instruction
 - Specifies multiple concurrent operations
 - ⇒ Very Long Instruction Word (VLIW)

322

3/2/2016 Chapter 4-The Processor- RST@NTHU



Scheduling Static Multiple Issue

- Compiler must remove some/all hazards
 - Reorder instructions into issue packets
 - No dependencies with a packet
 - Possibly some dependencies between packets
 - Varies between ISAs; compiler must know!
 - Pad with nop if necessary

3/2/2016 Chapter 4-The Processor- RST@NTHU



MIPS with Static Dual Issue

- Two-issue packets
 - One ALU/branch instruction
 - One load/store instruction
 - 64-bit aligned instruction
 - ALU/branch, then load/store
 - Pad an unused instruction with nop

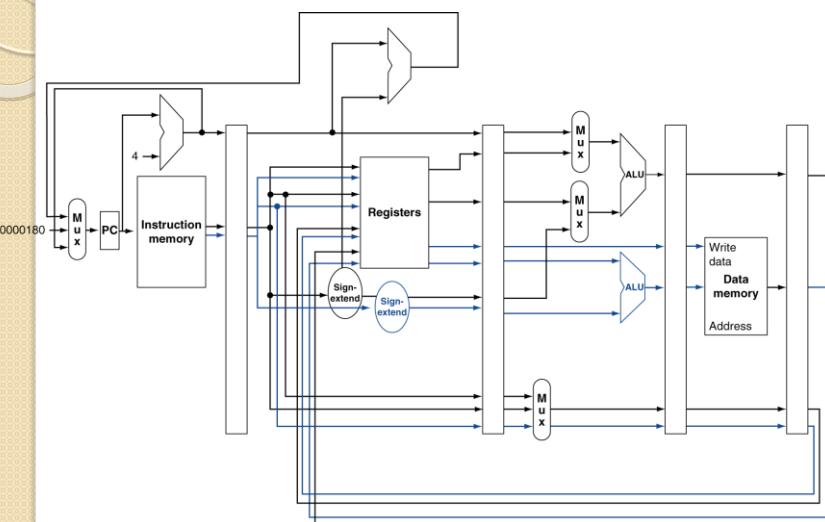
Address	Instruction type	Pipeline Stages							
		IF	ID	EX	MEM	WB			
n	ALU/branch								
n + 4	Load/store	IF	ID	EX	MEM	WB			
n + 8	ALU/branch		IF	ID	EX	MEM	WB		
n + 12	Load/store		IF	ID	EX	MEM	WB		
n + 16	ALU/branch			IF	ID	EX	MEM	WB	
n + 20	Load/store			IF	ID	EX	MEM	WB	

323

3/2/2016 Chapter 4-The Processor- RST@NTHU



MIPS with Static Dual Issue



324

3/2/2016 Chapter 4-The Processor- RST@NTHU



Hazards in the Dual-Issue MIPS

- More instructions executing in parallel
- EX data hazard
 - Forwarding avoided stalls with single-issue
 - Now can't use ALU result in load/store in same packet
 - add \$t0, \$s0, \$s1
load \$s2, 0(\$t0)
 - Split into two packets, effectively a stall
- Load-use hazard
 - Still one cycle use latency, but now two instructions
- More aggressive scheduling required

3/2/2016 Chapter 4-The Processor- RST@NTHU



Scheduling Example

- Schedule this for dual-issue MIPS

```
Loop: lw    $t0, 0($s1)      # $t0=array element
      addu $t0, $t0, $s2      # add scalar in $s2
      sw    $t0, 0($s1)      # store result
      addi $s1, $s1,-4        # decrement pointer
      bne  $s1, $zero, Loop   # branch $s1!=0
```

	ALU/branch	Load/store	cycle
Loop:	nop	lw \$t0, 0(\$s1)	1
	addi \$s1, \$s1,-4	nop	2
	addu \$t0, \$t0, \$s2	nop	3
	bne \$s1, \$zero, Loop	sw \$t0, 4(\$s1)	4

- IPC = 5/4 = 1.25 (c.f. peak IPC = 2)

325

3/2/2016 Chapter 4-The Processor- RST@NTHU



Loop Unrolling

- Replicate loop body to expose more parallelism
 - Reduces loop-control overhead
- Use different registers per replication
 - Called “register renaming”
 - Avoid loop-carried “anti-dependencies”
 - Store followed by a load of the same register
 - Aka “name dependence”
 - Reuse of a register name

326

3/2/2016 Chapter 4-The Processor- RST@NTHU



Loop Unrolling Example

	ALU/branch	Load/store	cycle
Loop:	addi \$s1, \$s1,-16	lw \$t0, 0(\$s1)	1
	nop	lw \$t1, 12(\$s1)	2
	addu \$t0, \$t0, \$s2	lw \$t2, 8(\$s1)	3
	addu \$t1, \$t1, \$s2	lw \$t3, 4(\$s1)	4
	addu \$t2, \$t2, \$s2	sw \$t0, 16(\$s1)	5
	addu \$t3, \$t4, \$s2	sw \$t1, 12(\$s1)	6
	nop	sw \$t2, 8(\$s1)	7
	bne \$s1, \$zero, Loop	sw \$t3, 4(\$s1)	8

- IPC = 14/8 = 1.75
 - Closer to 2, but at cost of registers and code size

326

3/2/2016 Chapter 4-The Processor- RST@NTHU



Dynamic Multiple Issue

- “Superscalar” processors
- CPU decides whether to issue 0, 1, 2, ... each cycle
 - Avoiding structural and data hazards
- Avoids the need for compiler scheduling
 - Though it may still help
 - Code semantics ensured by the CPU

327

3/2/2016 Chapter 4-The Processor- RST@NTHU



Dynamic Pipeline Scheduling

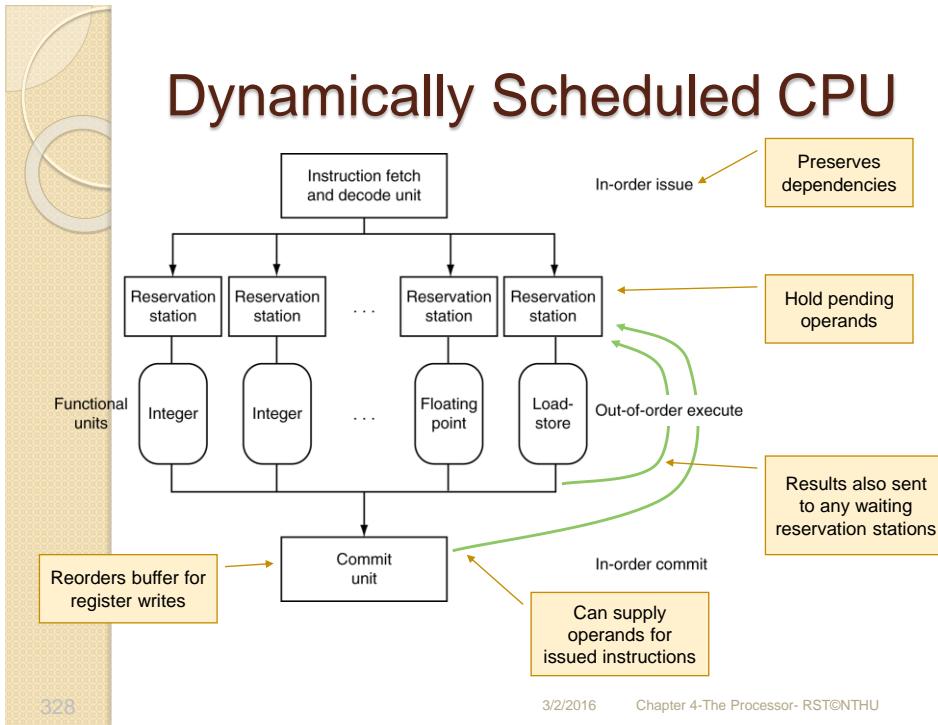
- Allow the CPU to execute instructions **out of order** to avoid stalls
 - But commit result to registers in order
- Example

```
lw      $t0, 20($s2)
addu   $t1, $t0, $t2
sub    $s4, $s4, $t3
slti   $t5, $s4, 20
```

 - Can start sub while addu is waiting for lw

327

3/2/2016 Chapter 4-The Processor- RST@NTHU



3/2/2016 Chapter 4-The Processor- RST@NTHU

Register Renaming

- Reservation stations and reorder buffer effectively provide register renaming
- On instruction issue to reservation station
 - If operand is available in register file or reorder buffer
 - Copied to reservation station
 - No longer required in the register; can be overwritten
 - If operand is not yet available
 - It will be provided to the reservation station by a function unit
 - Register update may not be required



400

3/2/2016 Chapter 4-The Processor- RST@NTHU

Speculation

- Predict branch and continue issuing
 - Don't commit until branch outcome determined
- Load speculation
 - Avoid load and cache miss delay
 - Predict the effective address
 - Predict loaded value
 - Load before completing outstanding stores
 - Bypass stored values to load unit
 - Don't commit load until speculation cleared

3/2/2016 Chapter 4-The Processor- RST@NTHU

Why Do Dynamic Scheduling?

- Why not just let the compiler schedule code?
- Not all stalls are predictable
 - e.g., cache misses
- Can't always schedule around branches
 - Branch outcome is dynamically determined
- Different implementations of an ISA have different latencies and hazards

3/2/2016 Chapter 4-The Processor- RST@NTHU



Does Multiple Issue Work?

The BIG Picture

- Yes, but **not** as much as we'd like
- Programs have real dependencies that limit ILP
- Some dependencies are hard to eliminate
 - e.g., pointer aliasing
- Some parallelism is hard to expose
 - Limited window size during instruction issue
- Memory delays and limited bandwidth
 - Hard to keep pipelines full
- Speculation can help if done well

3/2/2016 Chapter 4-The Processor- RST@NTHU



Power Efficiency

- Complexity of dynamic scheduling and speculations requires power
- Multiple simpler cores may be better

Microprocessor	Year	Clock Rate	Pipeline Stages	Issue width	Out-of-order/Speculation	Cores	Power
i486	1989	25MHz	5	1	No	1	5W
Pentium	1993	66MHz	5	2	No	1	10W
Pentium Pro	1997	200MHz	10	3	Yes	1	29W
P4 Willamette	2001	2000MHz	22	3	Yes	1	75W
P4 Prescott	2004	3600MHz	31	3	Yes	1	103W
Core	2006	2930MHz	14	4	Yes	2	75W
UltraSparc III	2003	1950MHz	14	4	No	1	90W
UltraSparc T1	2005	1200MHz	6	1	No	8	70W

3/2/2016 Chapter 4-The Processor- RST@NTHU

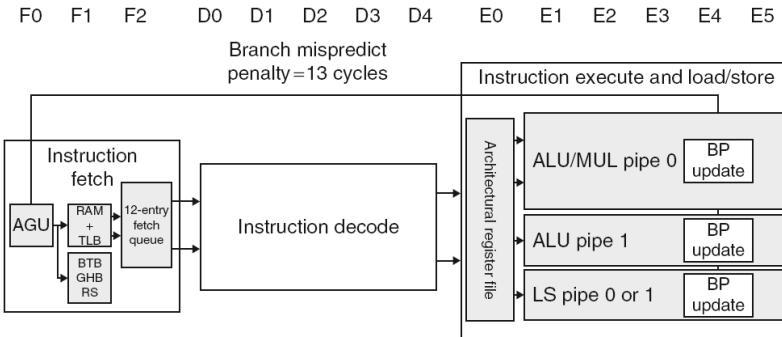
§4.11 Real Stuff: The ARM Cortex-A8 and Intel Core i7 Pipelines

Cortex A8 and Intel i7

Processor	ARM A8	Intel Core i7 920
Market	Personal Mobile Device	Server, cloud
Thermal design power	2 Watts	130 Watts
Clock rate	1 GHz	2.66 GHz
Cores/Chip	1	4
Floating point?	No	Yes
Multiple issue?	Dynamic	Dynamic
Peak instructions/clock cycle	2	4
Pipeline stages	14	14
Pipeline schedule	Static in-order	Dynamic out-of-order with speculation
Branch prediction	2-level	2-level
1 st level caches/core	32 KiB I, 32 KiB D	32 KiB I, 32 KiB D
2 nd level caches/core	128-1024 KiB	256 KiB
3 rd level caches (shared)	-	2- 8 MB

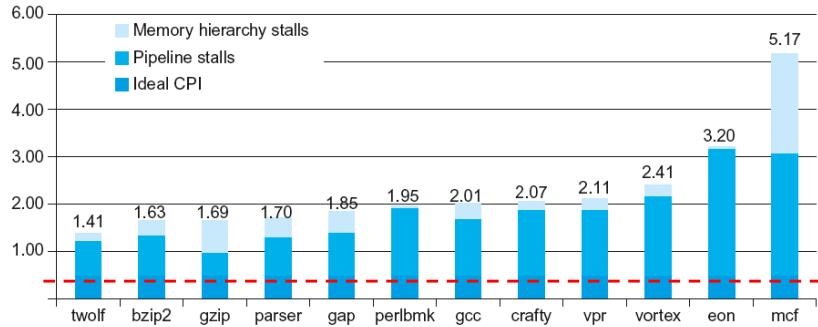
3/2/2016 Chapter 4-The Processor- RST@NTHU

ARM Cortex-A8 Pipeline



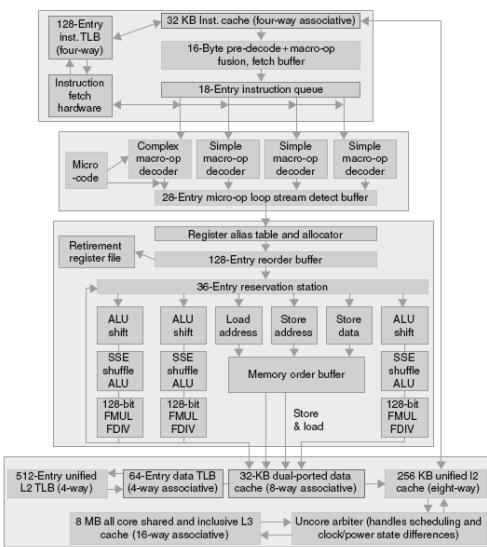
3/2/2016 Chapter 4-The Processor- RST@NTHU

ARM Cortex-A8 Performance



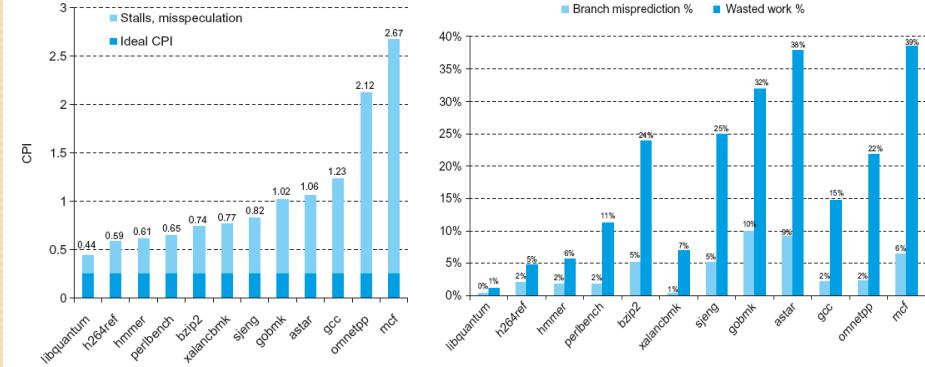
3/2/2016 Chapter 4-The Processor- RST@NTHU

Core i7 Pipeline



3/2/2016 Chapter 4-The Processor- RST@NTHU

Core i7 Performance



3/2/2016 Chapter 4-The Processor- RST@NTHU

§4.12 Instruction-Level Parallelism and Matrix Multiply

Matrix Multiply

Unrolled C code

```

1 #include <x86intrin.h>
2 #define UNROLL (4)
3
4 void dgemm ( int n, double* A, double* B, double* C)
5 {
6     for ( int i = 0; i < n; i+=UNROLL*4 )
7         for ( int j = 0; j < n; j++ ) {
8             __m256d c[4];
9             for ( int x = 0; x < UNROLL; x++ )
10                 c[x] = _mm256_load_pd(C+i+x*4+j*n);
11
12            for( int k = 0; k < n; k++ )
13            {
14                __m256d b = _mm256_broadcast_sd(B+k+j*n);
15                for (int x = 0; x < UNROLL; x++)
16                    c[x] = _mm256_add_pd(c[x],
17                                         _mm256_mul_pd(_mm256_load_pd(A+n*k+x*4+i), b));
18            }
19
20            for ( int x = 0; x < UNROLL; x++ )
21                _mm256_store_pd(C+i+x*4+j*n, c[x]);
22        }
23    }

```

3/2/2016 Chapter 4-The Processor- RST@NTHU



Matrix Multiply

Assembly code:

```

1 vmovapd (%r11),%ymm4
2 mov %rbx,%rax
3 xor %ecx,%ecx
4 vmovapd 0x20(%r11),%ymm3
5 vmovapd 0x40(%r11),%ymm2
6 vmovapd 0x60(%r11),%ymm1
7 vbroadcastsd (%rcx,%r9,1),%ymm0
8 add $0x8,%rcx # register %rcx = %rcx + 8
9 vmulpd (%rax),%ymm0,%ymm5
10 vaddpd %ymm5,%ymm4,%ymm4
11 vmulpd 0x20(%rax),%ymm0,%ymm5
12 vaddpd %ymm5,%ymm3,%ymm3
13 vmulpd 0x40(%rax),%ymm0,%ymm5
14 vmulpd 0x60(%rax),%ymm0,%ymm0
15 add %r8,%rax
16 cmp %r10,%rcx
17 vaddpd %ymm5,%ymm2,%ymm2
18 vaddpd %ymm0,%ymm1,%ymm1
19 jne 68 <dgemm+0x68>
20 add $0x1,%esi
21 vmovapd %ymm4,(%r11)
22 vmovapd %ymm3,0x20(%r11)
23 vmovapd %ymm2,0x40(%r11)
24 vmovapd %ymm1,0x60(%r11)

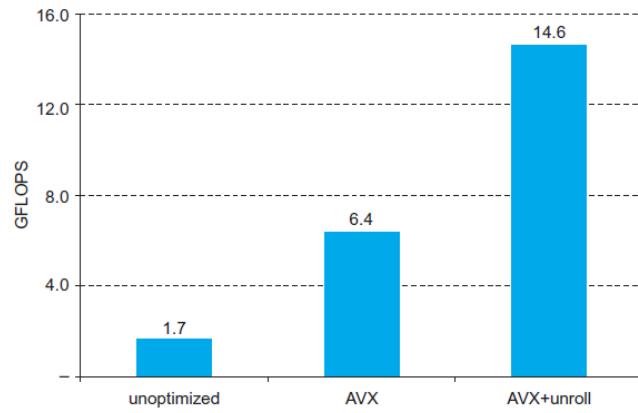
# Load 4 elements of C into %ymm4
# register %rax = %rbx
# register %ecx = 0
# Load 4 elements of C into %ymm3
# Load 4 elements of C into %ymm2
# Load 4 elements of C into %ymm1
# Make 4 copies of B element
# Parallel mul %ymm1,4 A elements
# Parallel add %ymm5, %ymm4
# Parallel mul %ymm1,4 A elements
# Parallel add %ymm5, %ymm3
# Parallel mul %ymm1,4 A elements
# Parallel mul %ymm1,4 A elements
# register %rax = %rax + %r8
# compare %r8 to %rax
# Parallel add %ymm5, %ymm2
# Parallel add %ymm0, %ymm1
# jump if not %r8 != %rax
# register %esi = %esi + 1
# Store %ymm4 into 4 C elements
# Store %ymm3 into 4 C elements
# Store %ymm2 into 4 C elements
# Store %ymm1 into 4 C elements

```

3/2/2016 Chapter 4-The Processor- RST@NTHU



Performance Impact



3/2/2016 Chapter 4-The Processor- RST@NTHU

Fallacies

- Pipelining is easy (!)
 - The basic idea is easy
 - The devil is in the details
 - e.g., detecting data hazards
- Pipelining is independent of technology
 - So why haven't we always done pipelining?
 - More transistors make more advanced techniques feasible
 - Pipeline-related ISA design needs to consider technology trends
 - e.g., predicated instructions

407

3/2/2016 Chapter 4-The Processor- RST©NTHU

Pitfalls

- Poor ISA design can make pipelining harder
 - e.g., complex instruction sets (VAX, IA-32)
 - Significant overhead to make pipelining work
 - IA-32 micro-op approach
 - e.g., complex addressing modes
 - Register update side effects, memory indirection
 - e.g., delayed branches
 - Advanced pipelines have long delay slots

407

3/2/2016 Chapter 4-The Processor- RST©NTHU

Concluding Remarks

- ISA influences design of datapath and control
- Datapath and control influence design of ISA
- Pipelining improves instruction throughput using parallelism
 - More instructions completed per second
 - Latency for each instruction not reduced
- Hazards: structural, data, control
- Multiple issue and dynamic scheduling (ILP)
 - Dependencies limit achievable parallelism
 - Complexity leads to the power wall